

Machine Intelligence:: Deep Learning

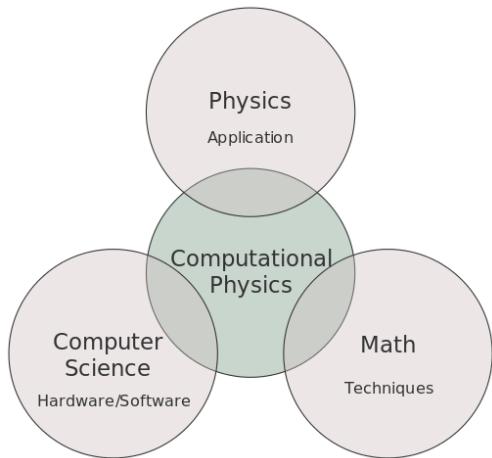
Week 1

Beate Sick, Oliver Dürr, Jonas Brändli

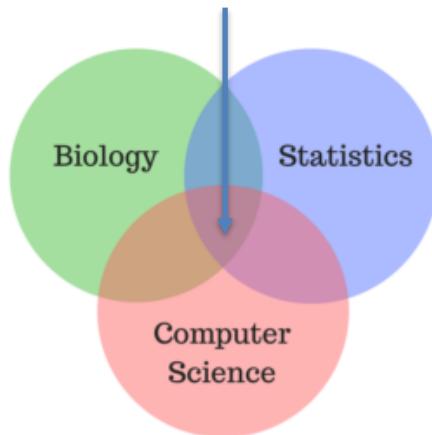
Institut für Datenanalyse und Prozessdesign
Zürcher Hochschule für Angewandte Wissenschaften

Oliver's Background

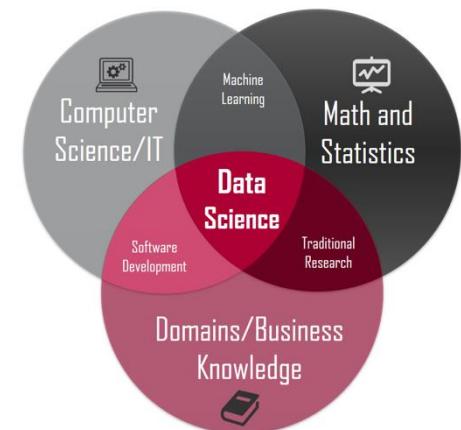
Computational Physics



Bioinformatics

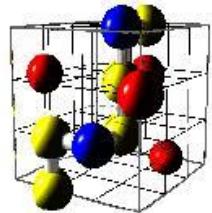


Data Science



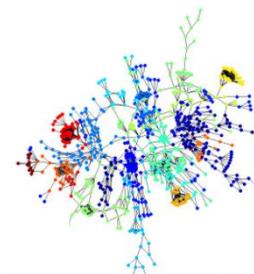
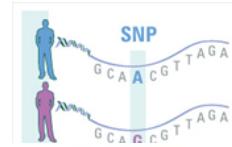
1990's

Uni-Konstanz

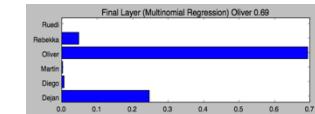
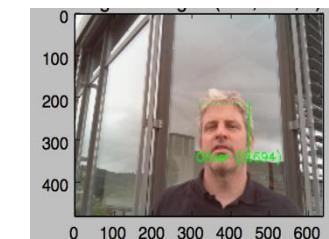


2000's

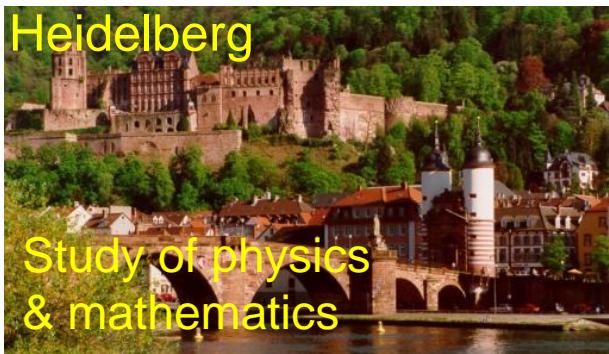
Genedata Basel



eclipse



Beate's Background



Head of bioinformatics
Focus: Gene expression

Researcher and Professor for applied statistics

Focus: deep learning

Researcher and lecturer
Focus: Biostatistics, DL

Tell us something about you

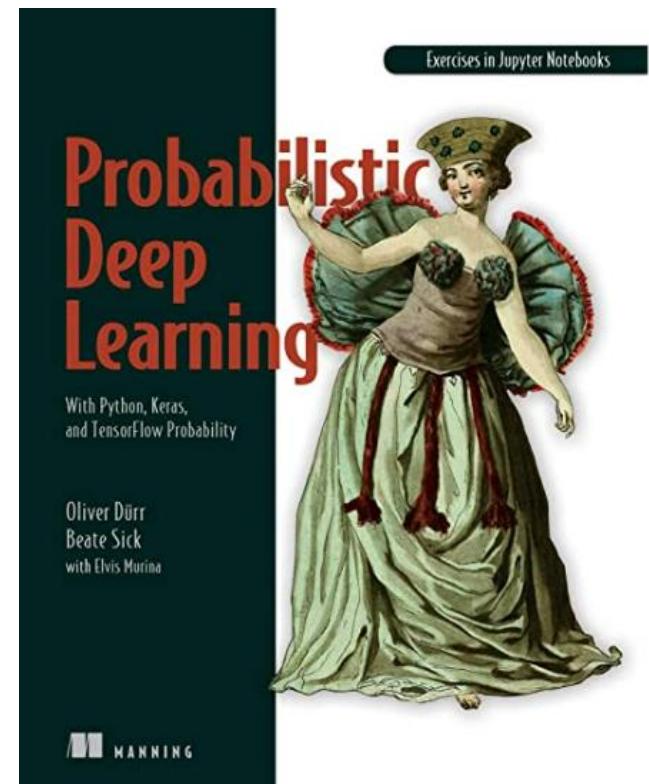
- Computer Science Background
 - Fluent in python?
- Statistics / Math
 - Who visited CAS StMo (statistisches Modellieren)?
 - What is a distribution?
 - Vector times Matrix?
 - Please make sure to check
https://tensorchiefs.github.io/dl_course_2022/prerequisites.html
- Any contacts with deep learning yet?

Technical details for this course

- Running the code:
 - Colab Notebooks
 - needs no installation, only internet and google account
 - Anaconda
 - Installation by your own

Material for the course

- Website and Github repository
 - The CAS Deep Learning Course
 - https://tensorchiefs.github.io/dl_course_2022/
- Our Book “Probabilistic Deep Learning”
 - Can be used in addition to the course
 - https://www.manning.com/books/probabilistic-deep-learning-with-python?a_aid=probabilistic_deep_learning&a_bid=78e55885
 - https://github.com/tensorchiefs/dl_book



Organizational Issues: ~~Test~~ Projects

- Projects (2-3 People)
- Presented on the last day
 - Spotlight talk (5 Minutes)
 - Poster
- Topics
 - You can / should choose a topic of your own (please discuss your topic with us by week4 latest)
 - Possible Topics (see website)
 - Take part in a Kaggle Competition (e.g. Leaf Classification / Dogs vs. Cats)
 - Music classification
 - Polar bear detection
 - ...
- Computing: colab, laptop (or cloud computing)

Organizational Issues: Times

- Dates and times: see our webpage
- Afternoon sessions
 - 13:30-17:00
- Theory and exercises will be mixed
 - Could be 50 minutes theory 30 minutes exercises
 - Could be vice versa
- **Please interrupt us if something is unclear! The less we talk the better!**

Outline of the DL Module (tentative)

- Day 1: Jumpstart to DL
 - What is DL
 - Basic Building Blocks
 - Keras
- Day 2: CNN I
 - ImageData
- Day 3: CNN II and RNN
 - Tips and Tricks
 - Modern Architectures
 - 1-D Sequential Data
- Day 4: Looking at details
 - Linear Regression
 - Backpropagation
 - Resnet
 - Likelihood principle
- Day 5: Probabilistic Aspects
 - TensorFlow Probability (TFP)
 - Negative Loss Likelihood NLL
 - Count Data
- Day 6: Probabilistic models in the wild
 - Complex Distributions
 - Generative modes with normalizing flows
- Day 7: Uncertainty in DL
 - Bayesian Modeling
- Day 8: Uncertainty cont'd
 - Bayesian Neural Networks
 - Projects

Day 1-4 should get you ready for your project.

Learning Objectives for today

- Get a rough idea what the DL is about
- Get a first idea on patterns in NN / DL
 - Flow of tensors
 - Matrix and Tensor operations
 - Backpropagation
 - To fit the weights of a network efficiently
- Framework
 - Introduction to Keras

Introduction to Deep Learning --what's the hype about?

Machine Perception

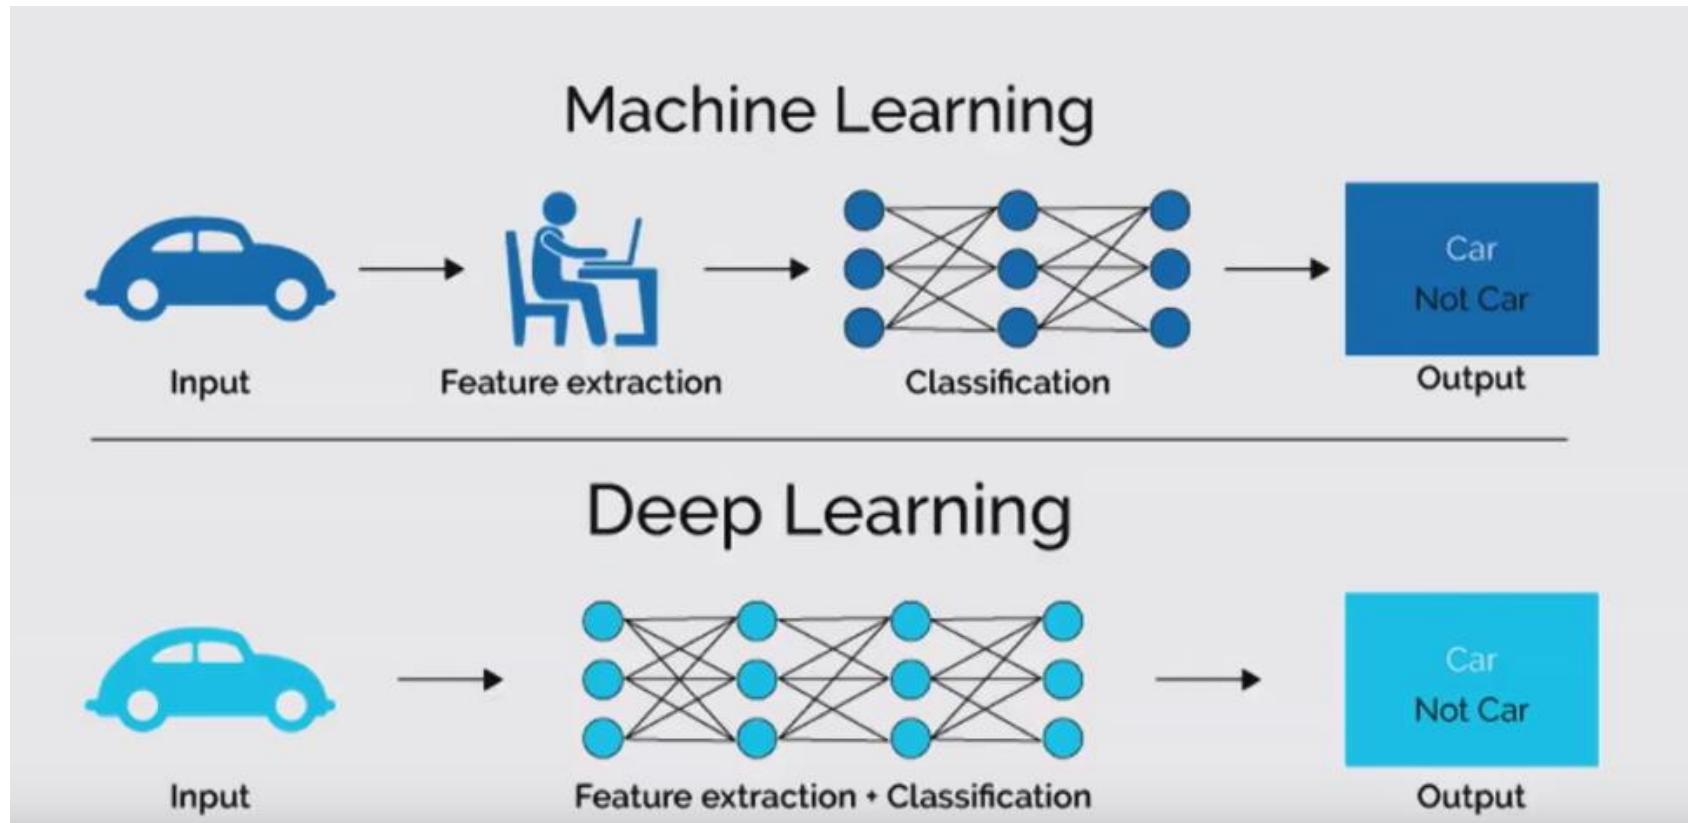
- Computers have been quite bad in things which are easy for humans (images, text, sound)
- A Kaggle contest 2012
- In the following we explain why

Kaggle dog vs cat competition



Deep Blue beat Kasparov at chess in 1997.
Watson beat the brightest trivia minds at Jeopardy in 2011.
Can you tell Fido from Mittens in 2013?

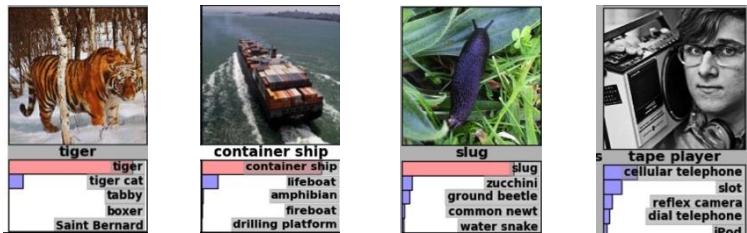
Deep Learning vs. Machine Learning



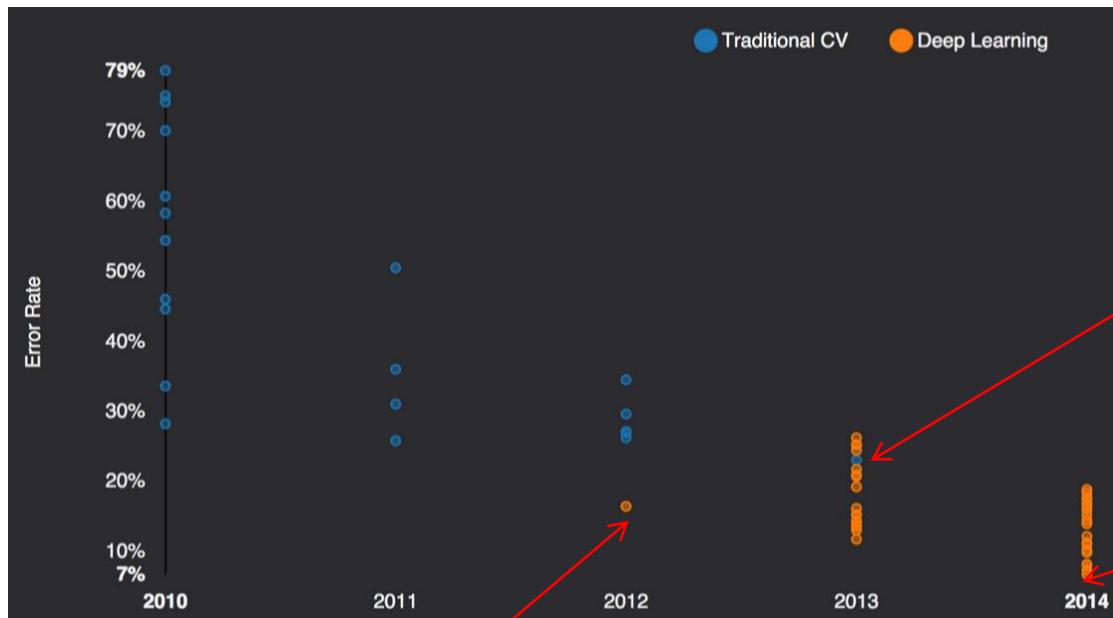
The most convincing case for
DL (subjective view)

Why DL: Imagenet 2012, 2013, 2014, 2015

1000 classes
1 Mio samples



...



Human: 5% misclassification

Only one non-CNN approach in 2013

GoogLeNet 6.7%

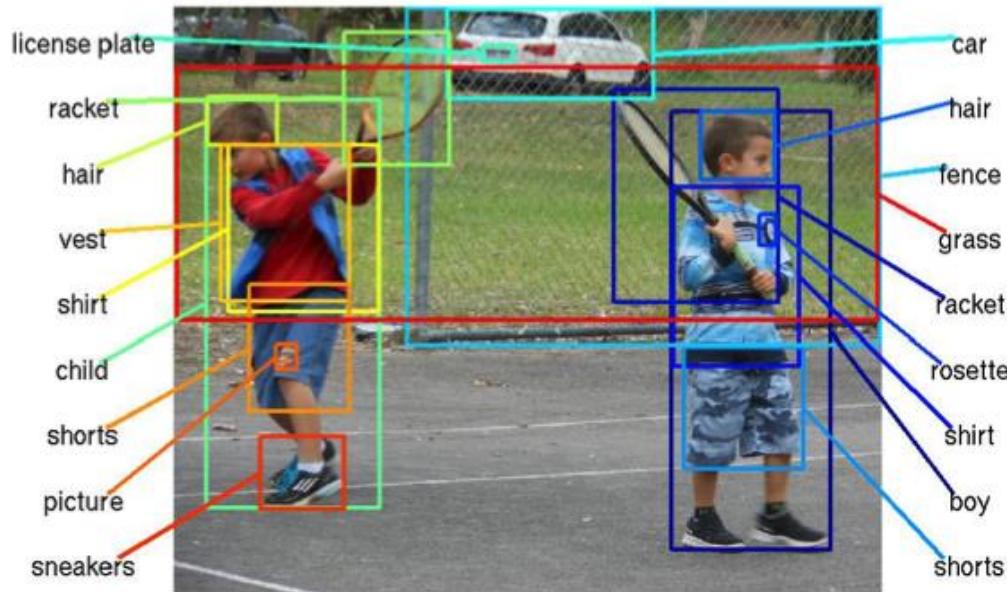
A. Krizhevsky
first CNN in 2012
Und es hat zoom gemacht

2015: It gets tougher

- 4.95% Microsoft ([Feb 6](#) surpassing human performance 5.1%)
- 4.8% Google ([Feb 11](#)) -> further improved to 3.6 (Dec)?
- 4.58% Baidu (May 11 [banned due to too many submissions](#))
- 3.57% Microsoft (Resnet winner 2015) → task solved!

The computer vision success story

- With DL it took approx. 3 years to solve object detection and other computer vision task



Images from cs229n

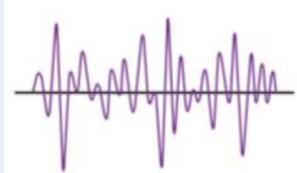


Deep Blue beat Kasparov at chess in 1997.
Watson beat the brightest trivia minds at Jeopardy in 2011.
Can you tell Fido from Mittens in 2013?



"man in black shirt is playing guitar."

Use cases of deep learning

Input x to DL model	Output y of DL model	Application
Images 	Label “Tiger”	Image classification
Audio 	Sequence / Text “see you tomorrow”	Voice Recognition
Sequence (prompt) An astronaut riding a horse in a photorealistic style		Image Generation
Sequences (prompt) “Hallo, wie?”	Next word “geht”	Language Models
Simple number (age) age=52	Simple number (SPB) sbp = 152	Simple Regression Educational

Deep Learning öffnet Tür zu hören, sehen und Texten.

Status Quo: kein Verstehen aber Erfassung statistische Zusammenhänge.

This is the new shit: ChatGPT



Die gefühlte Revolution

4. Dezember 2022, 18:51 Uhr | Lesezeit: 3 min



Das kommt heraus, wenn man der künstlichen Intelligenz Dall-E die Anweisung gibt: "Ein Roboter lässt beim Turing-Test einen Menschen glauben, dass sie ein Mensch ist, im Stil von Kehinde Wiley." (Foto: Dall-E-Bild: SZ)



GPT (short for "Generative Pre-training Transformer")

is a type of language processing AI model developed by OpenAI. It is a large, deep learning model that has been trained on a diverse range of texts and can generate human-like text when given a prompt.

Deep learning Artificial Intelligence?

All the impressive achievements of deep learning amount to just curve fitting

Juda Pearl, 2018

Pearl's ladder of causality

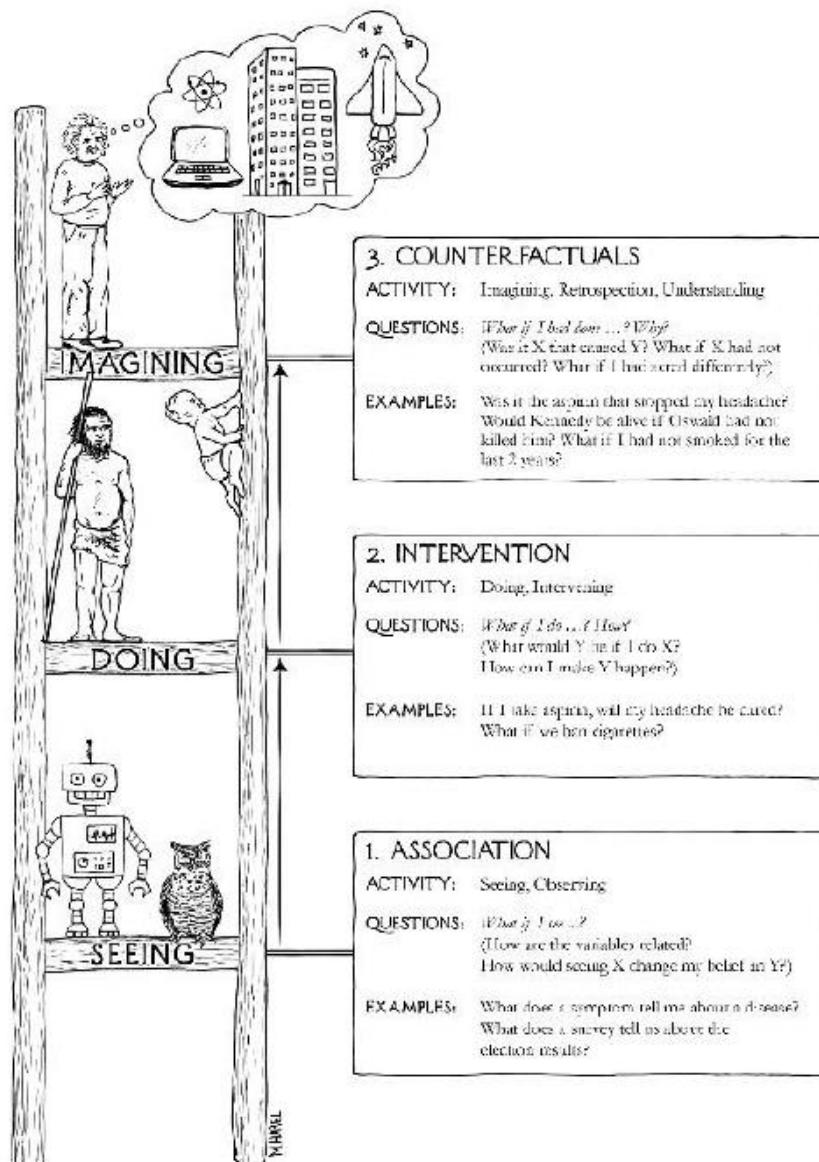
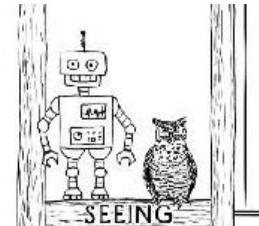


FIGURE 1.2. The Ladder of Causation, with representative organisms at each level. Most animals, as well as present-day learning machines, are on the first

On the first rung of the ladder DL is currently as good as a ensemble of pigeons ;-)



<https://www.youtube.com/watch?v=NsV6S8EsC0E>



OPEN ACCESS

Citation: Levenson RM, Krupinski EA, Navarro VM, Wasserman EA (2015) Pigeons (*Columba livia*) as Trainable Observers of Pathology and Radiology Breast Cancer Images. PLoS ONE 10(11): e0141357. doi:10.1371/journal.pone.0141357

Editor: Jonathan A Coles, Glasgow University, UNITED KINGDOM

Received: August 25, 2015

Accepted: October 7, 2015

Published: November 18, 2015

RESEARCH ARTICLE

Pigeons (*Columba livia*) as Trainable Observers of Pathology and Radiology Breast Cancer Images

Richard M. Levenson^{1*}, Elizabeth A. Krupinski³, Victor M. Navarro², Edward A. Wasserman^{2*}

1 Department of Pathology and Laboratory Medicine, University of California Davis Medical Center, Sacramento, California, United States of America, **2** Department of Psychological and Brain Sciences, The University of Iowa, Iowa City, Iowa, United States of America, **3** Department of Radiology & Imaging Sciences, College of Medicine, Emory University, Atlanta, Georgia, United States of America

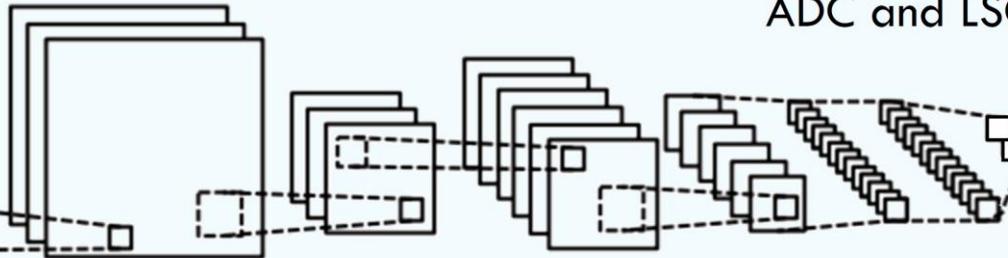
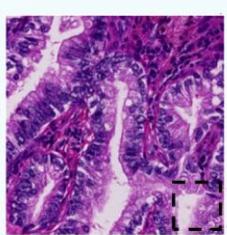
* levenson@ucdavis.edu (RML); ed-wasserman@uiowa.edu (EAW)

Abstract

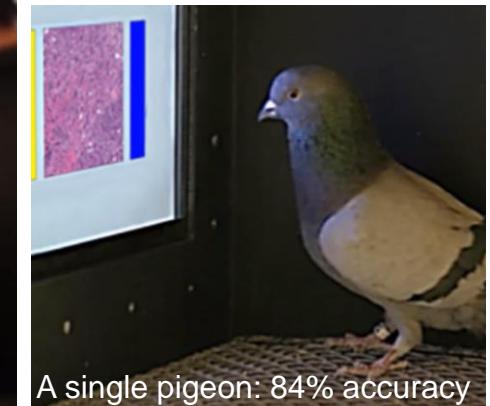
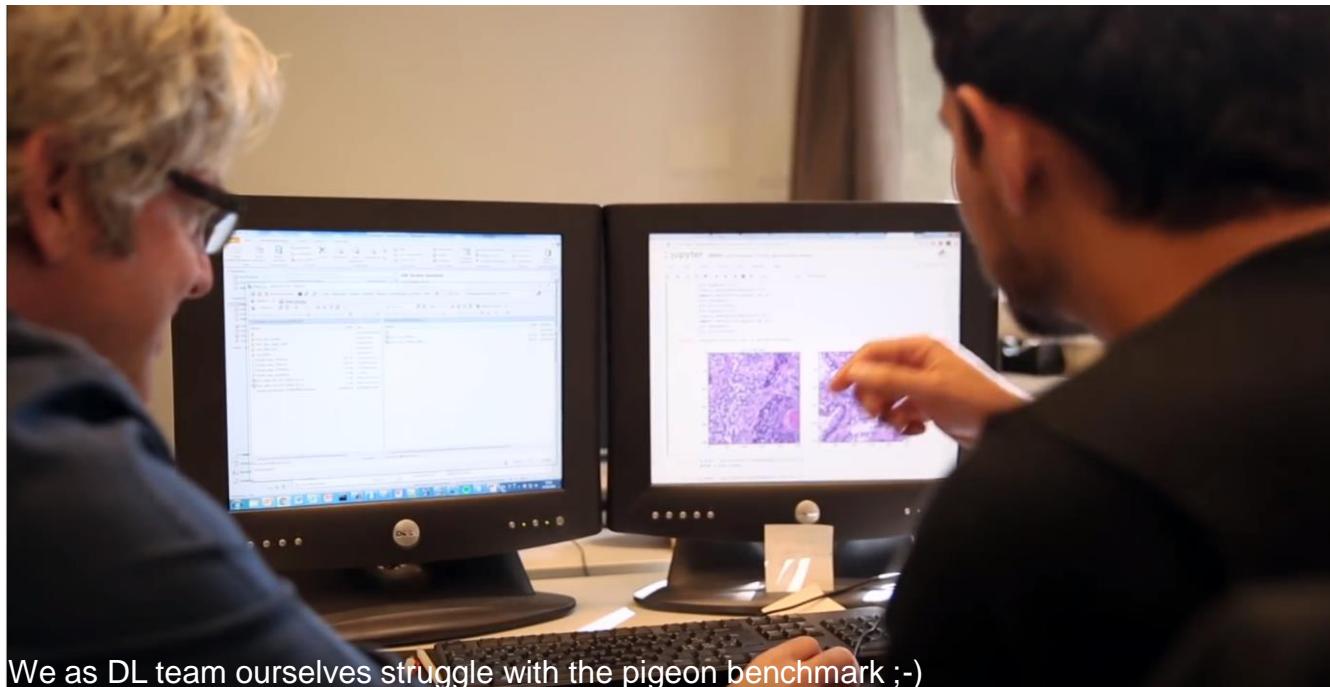
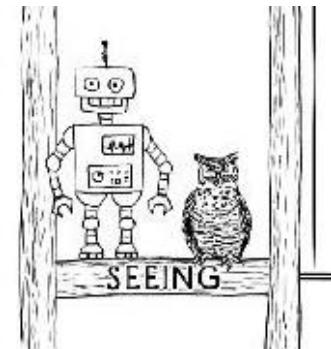
Pathologists and radiologists spend years acquiring and refining their medically essential visual skills, so it is of considerable interest to understand how this process actually unfolds and what image features and properties are critical for accurate diagnostic performance. Key insights into human behavioral tasks can often be obtained by using appropriate animal models. We report here that pigeons (*Columba livia*)—which share many visual system properties with humans—can serve as promising surrogate observers of medical images, a capability not previously documented. The birds proved to have a remarkable ability to distinguish benign from malignant human breast histopathology after training with differential food reinforcement; even more importantly, the pigeons were able to generalize what they had learned when confronted with novel image sets. The birds' histological accuracy, like that of humans, was modestly affected by the presence or absence of color as well as by degrees of image compression, but these impacts could be ameliorated with further training. Turning to radiology, the birds proved to be similarly capable of detecting cancer-relevant microcalcifications on mammogram images. However, when given a different (and for humans quite difficult) task—namely, classification of suspicious mammographic densities (masses)—the pigeons proved to be capable only of image memorization and were unable

On the first rung of the ladder DL is currently as good as an ensemble of pigeons

Our DL model achieves ~90% accuracy on image level



Probability for
ADC and LSCC



First Neural Network

The Single Cell: Biological Motivation

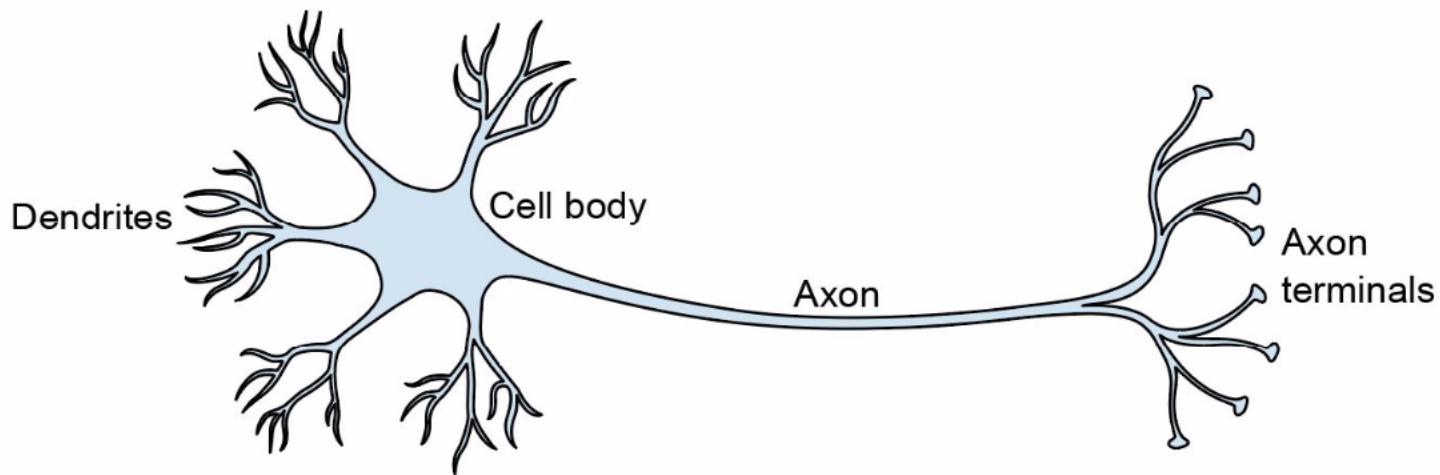


Figure 2.2 A single biological brain cell. The neuron receives the signal from other neurons via its dendrites shown on the left. If the cumulated signal exceeds a certain value, an impulse is sent via the axon to the axon terminals, which, in turn, couples to other neurons.

Neural networks are **loosely** inspired by how the brain works

The Single Cell: Mathematical Abstraction

$$z = b + x_1 \cdot w_1 + x_2 \cdot w_2 + \cdots x_p \cdot w_p$$

$$z = b + \sum x_i \cdot w_i = b + \mathbf{x} \cdot \mathbf{w}$$

Activation (many possibilities)

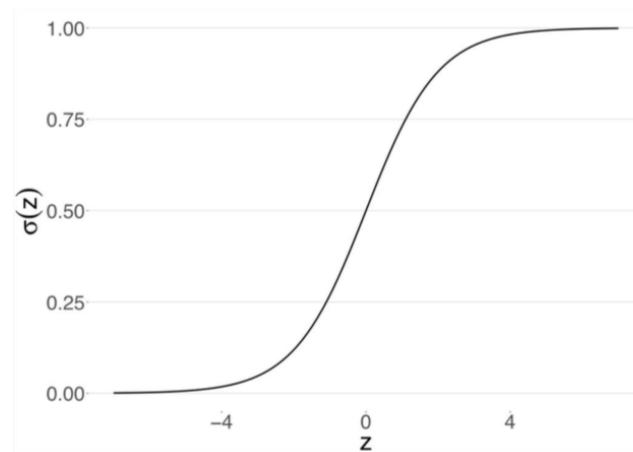
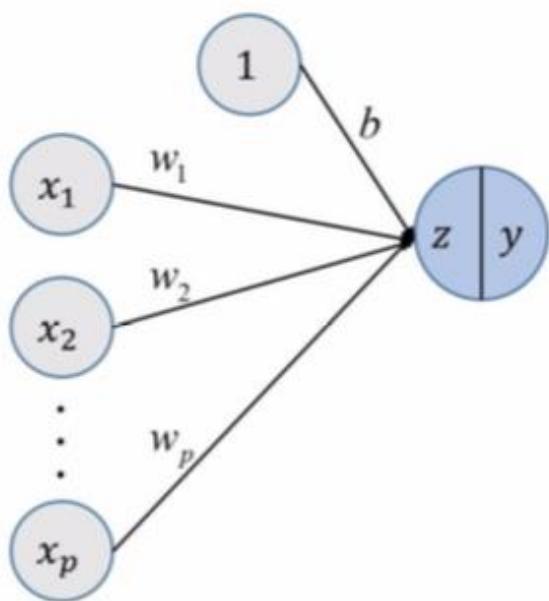


Figure 2.3 The mathematical abstraction of a brain cell (an artificial neuron). The value z is computed as the weighted sum of the p input values, x_1 to x_p , and a bias term b that shifts up or down the resulting weighted sum of the inputs. The value y is computed from z by applying an activation function.

```
# definition of the sigmoid function
def sigmoid(z):
    return (1 / (1 + np.exp(-z)))
```

Toy Task

- Task tell fake from real banknotes
- Banknotes described by two features

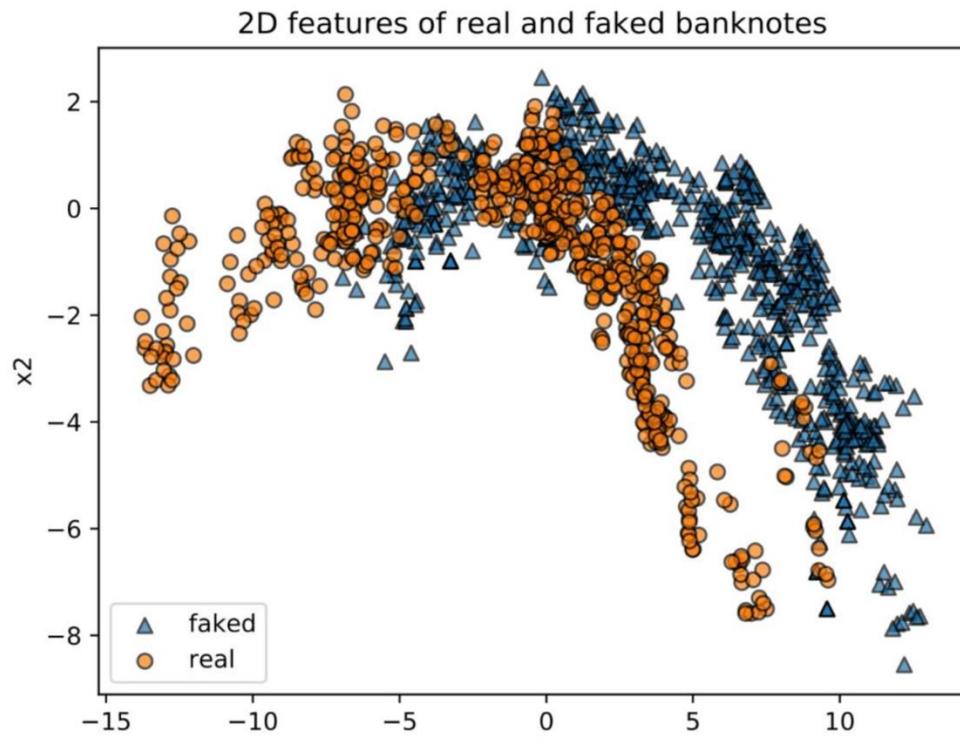
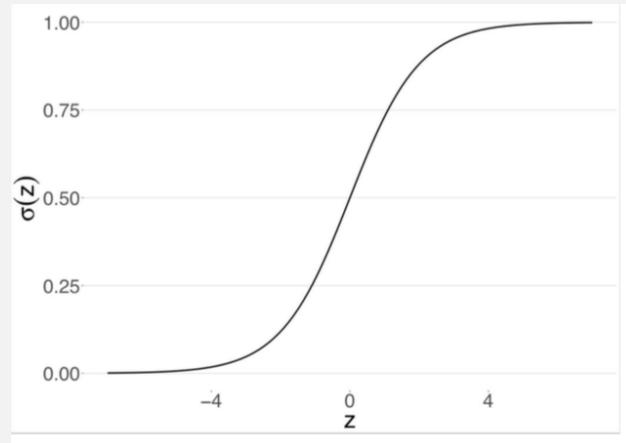
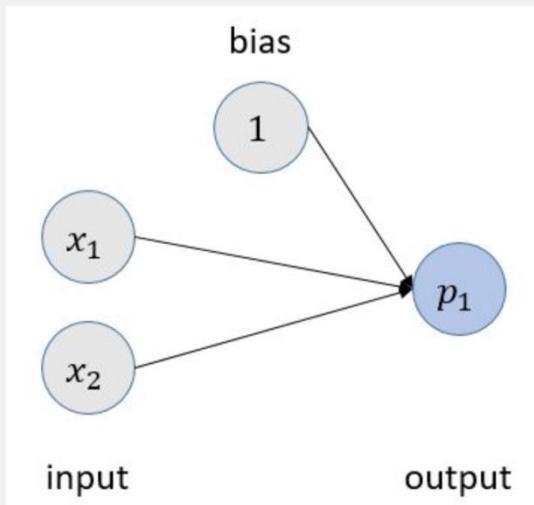


Figure 2.5 The (training) data points for the real and faked banknotes

Exercise: Part 1



Model: The above network models the **probability** p_1 that a given banknote is false.

TASK (with pen and paper)

The weights (determined by a training procedure later) are given by

$$w_1 = 0.3, w_2 = 0.1, \text{ and } b = 1.0$$

The probability can be calculated from z using the function $\text{sigmoid}(z)$

What is the probability that a banknote, that is characterized by $x_1=1$ and $x_2 = 2.2$, is a faked banknote?

GPUs love Vectors

 $F^{\mu\nu}$

In Math:

$$p_1 = \text{sigmoid} \left((x_1 \quad x_2) \cdot \begin{pmatrix} w_1 \\ w_2 \end{pmatrix} + b \right)$$

In code:

```
## function to return the probability output after the matrix multiplication
def predict_no_hidden(X):
    return sigmoid(np.matmul(X,W)+b)
```

Recap: Matrix Multiplication aka dot-product of matrices

We can only multiply matrices if their dimensions are compatible.

$$\mathbf{A} \times \mathbf{B} = \mathbf{C}$$

$$(m \times n) \times (n \times p) = (m \times p)$$

$$\begin{array}{c} \mathbf{A}_{3 \times 3} \quad \times \quad \mathbf{B}_{3 \times 2} \quad = \quad \mathbf{C}_{3 \times 2} \\ \left[\begin{array}{ccc} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ \boxed{a_{31}} & a_{32} & a_{33} \end{array} \right] \times \left[\begin{array}{cc} b_{11} & b_{12} \\ b_{21} & b_{22} \\ \boxed{b_{31}} & b_{32} \end{array} \right] = \left[\begin{array}{cc} c_{11} & c_{12} \\ c_{21} & c_{22} \\ \boxed{c_{31}} & c_{32} \end{array} \right] \end{array}$$

$$c_{11} = a_{11}b_{11} + a_{12}b_{21} + a_{13}b_{31}$$

$$c_{12} = a_{11}b_{12} + a_{12}b_{22} + a_{13}b_{32}$$

$$c_{21} = a_{21}b_{11} + a_{22}b_{21} + a_{23}b_{31}$$

$$c_{22} = a_{21}b_{12} + a_{22}b_{22} + a_{23}b_{32}$$

$$\boxed{c_{31} = a_{31}b_{11} + a_{32}b_{21} + a_{33}b_{31}}$$

$$c_{32} = a_{31}b_{12} + a_{32}b_{22} + a_{33}b_{32}$$

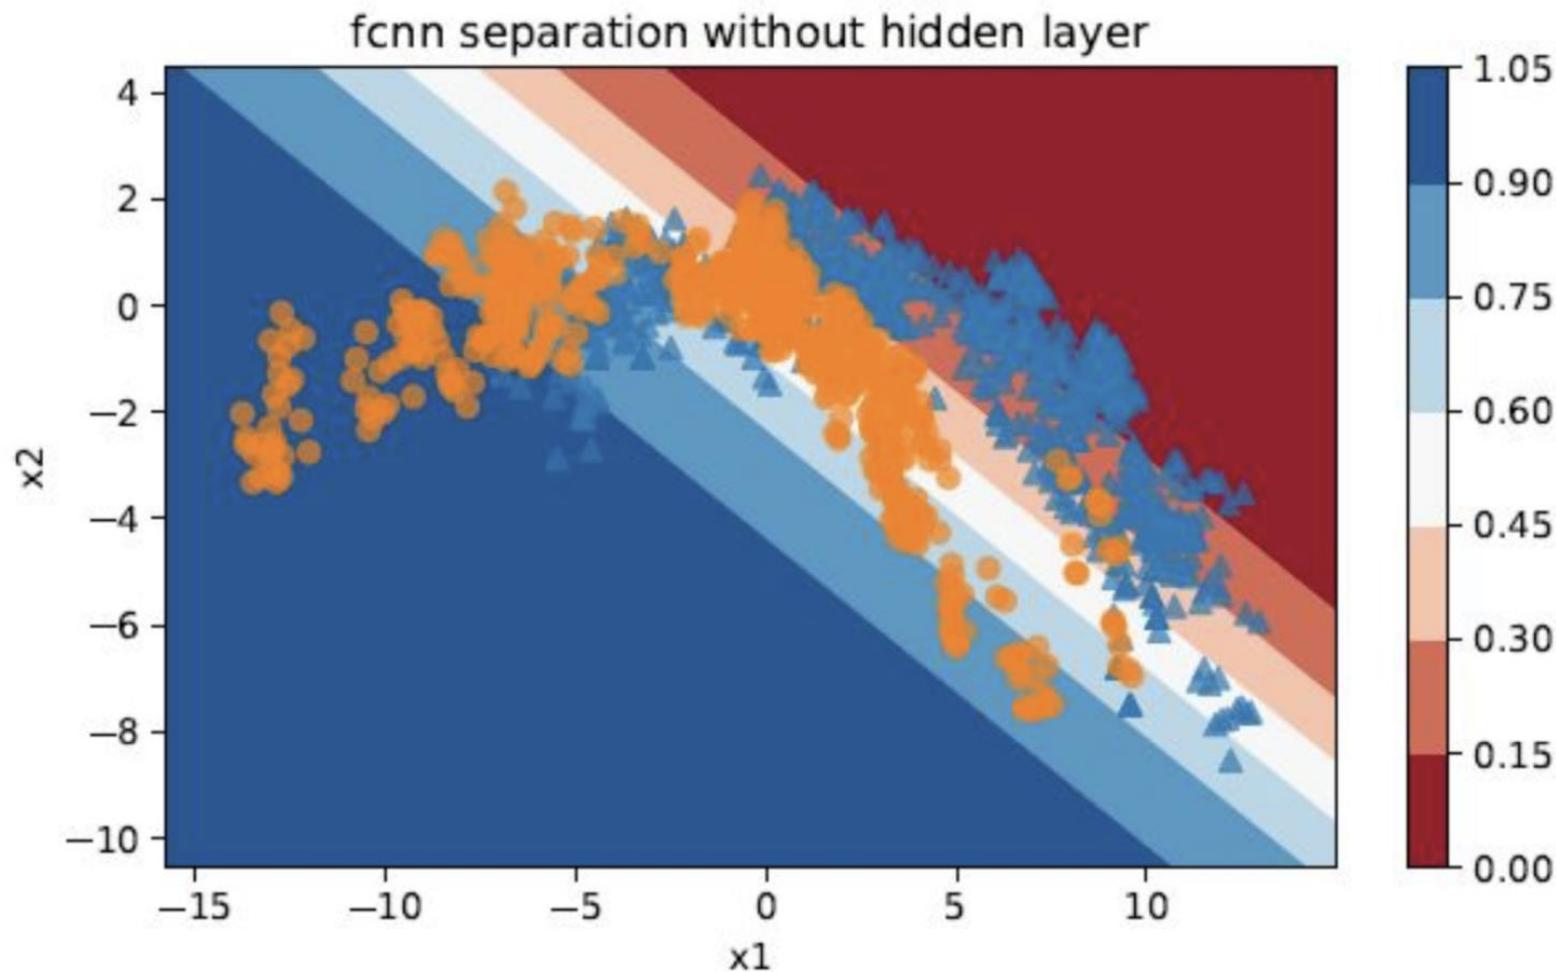
Example:

$$\mathbf{A}_{2 \times 2} = \begin{pmatrix} 2 & 1 \\ 0 & 3 \end{pmatrix}$$

$$\mathbf{B}_{2 \times 3} = \begin{pmatrix} 3 & \boxed{1} & 7 \\ 8 & \boxed{2} & 4 \end{pmatrix}$$

$$\mathbf{C}_{2 \times 3} = \mathbf{A}_{2 \times 2} \cdot \mathbf{B}_{2 \times 3} = \begin{pmatrix} 11 & 4 & 18 \\ 24 & \boxed{6} & 12 \end{pmatrix}$$

Result (see later in the notebook)



General rule: Networks without hidden layer have linear decision boundary.

To go deep non-linear activation functions are needed

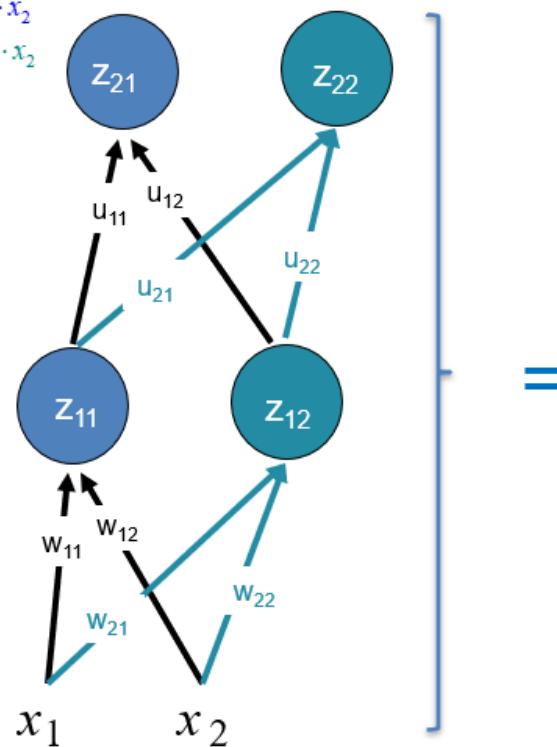
2 linear layers can be replaced by 1 linear layer -> can't go deep with linear layers!

$$z = (x \cdot W) \cdot U = x \cdot (W \cdot U) = x \cdot V$$

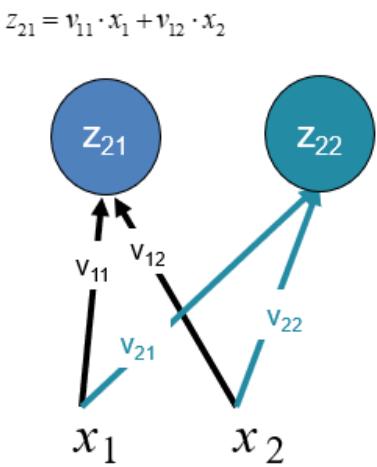
$$\begin{aligned} z_{21} &= z_{11} \cdot u_{11} + z_{12} \cdot u_{12} = (w_{11} \cdot x_1 + w_{12} \cdot x_2) \cdot u_{11} + (w_{21} \cdot x_1 + w_{22} \cdot x_2) \cdot u_{12} \\ &= x_1 \cdot (w_{11} \cdot u_{11} + w_{21} \cdot u_{12}) + x_2 \cdot (w_{12} \cdot u_{11} + w_{22} \cdot u_{12}) \end{aligned}$$

$$z_{11} = w_{11} \cdot x_1 + w_{12} \cdot x_2$$

$$z_{12} = w_{21} \cdot x_1 + w_{22} \cdot x_2$$



=



$$z_{21} = v_{11} \cdot x_1 + v_{12} \cdot x_2$$

$$v_{11} = w_{11} \cdot u_{11} + w_{21} \cdot u_{12}$$

$$v_{12} = w_{12} \cdot u_{11} + w_{22} \cdot u_{12}$$

$$v_{21} = w_{11} \cdot u_{21} + w_{21} \cdot u_{22}$$

$$v_{22} = w_{12} \cdot u_{21} + w_{22} \cdot u_{22}$$

Remark: biases are ignored here, but do not change fact

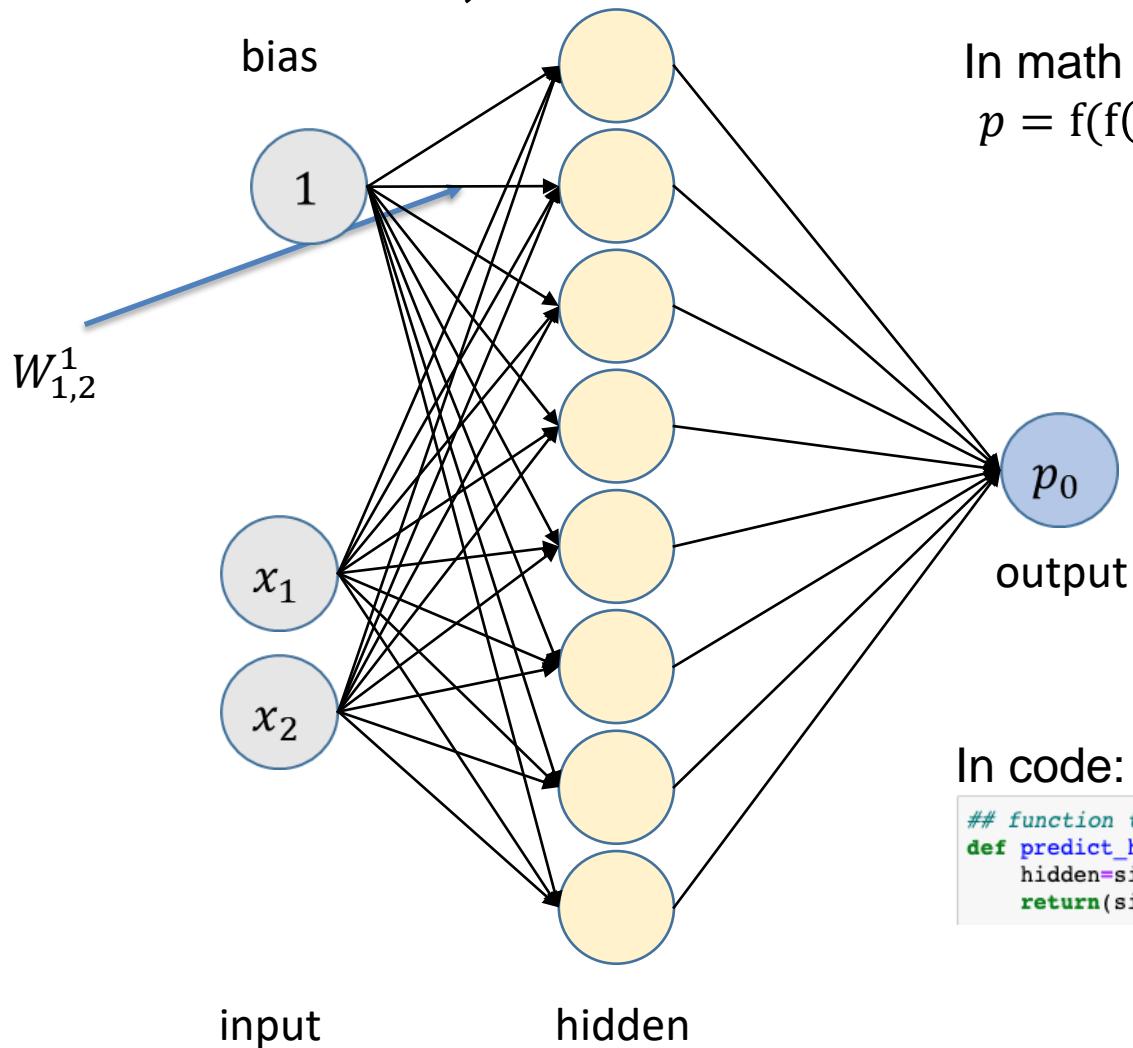
A close-up shot from the movie Inception. Two men in dark suits are facing each other. The man on the left has his eyes closed and a slight smile. The man on the right has his eyes open and is looking directly at the other man. The lighting is warm and dramatic.

WE NEED TO GO

DEEPER

A first deep network

$W_{\text{from,to}}$



In math ($f = \text{sigmoid}$)

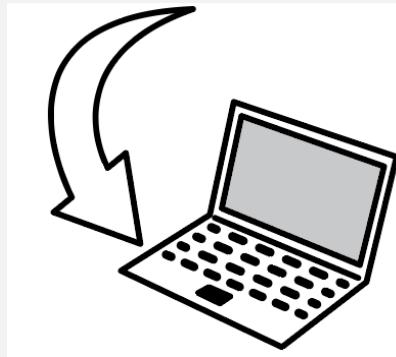
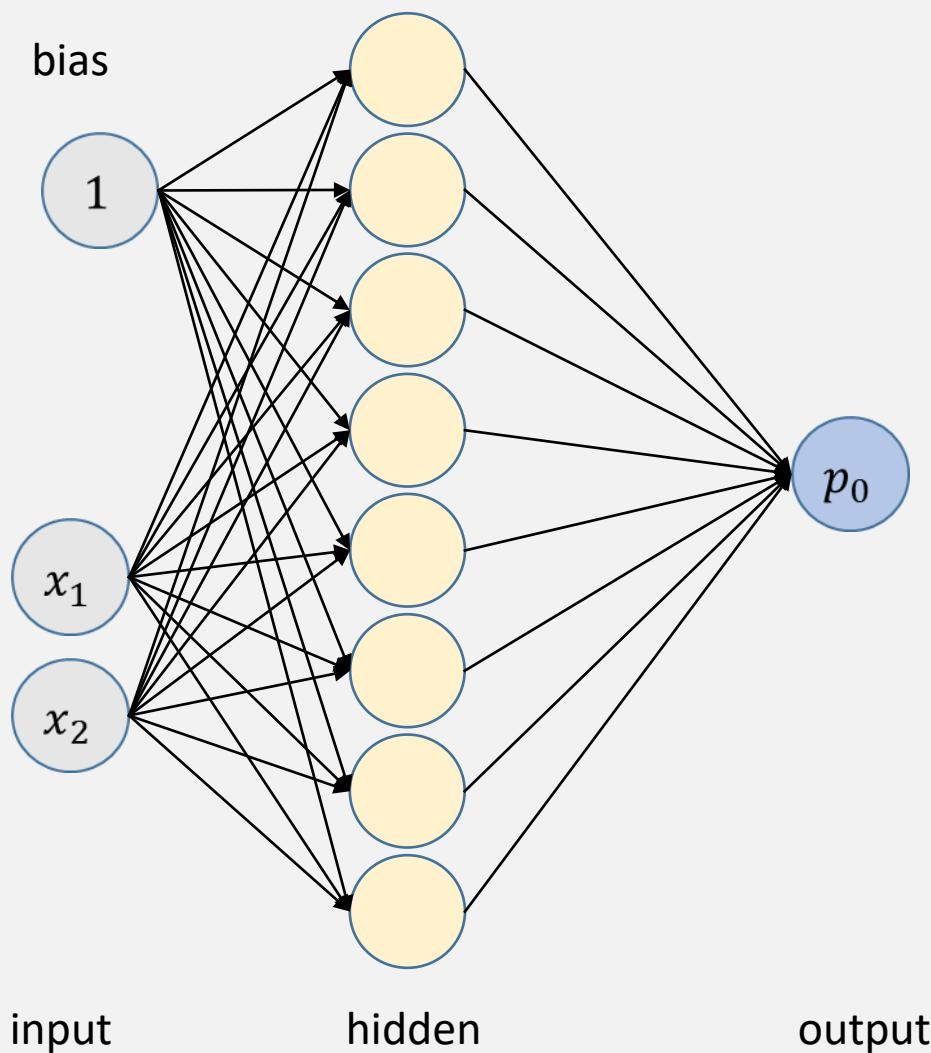
$$p = f(f(X \cdot W_1 + b_1) \cdot W_2 + b_2)$$

output

In code:

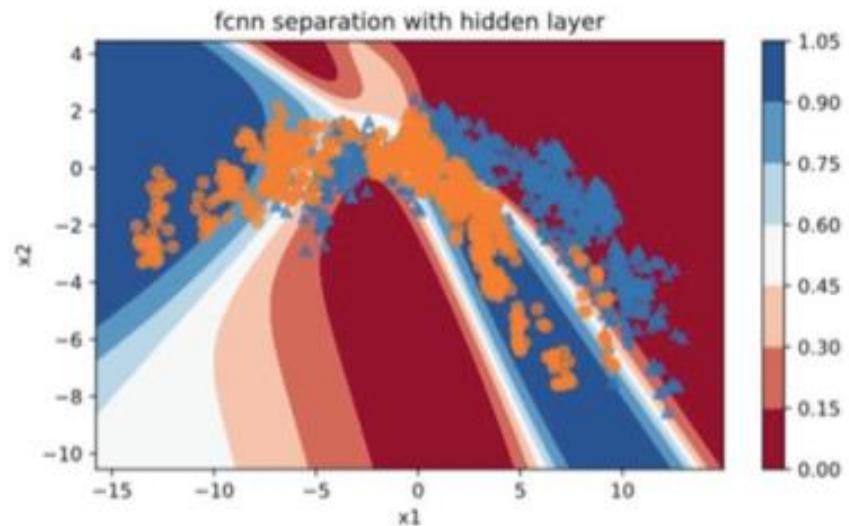
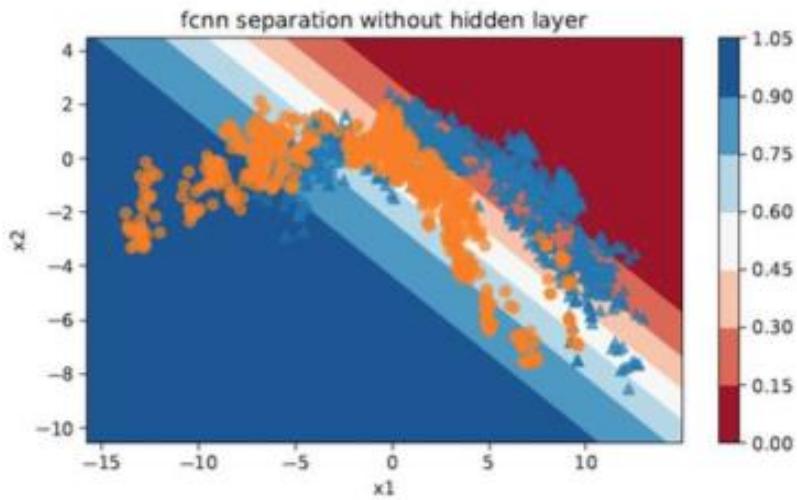
```
## function to return the probability output after the hidden layer
def predict_hidden(X):
    hidden=sigmoid(np.matmul(X,W1)+b1)
    return(sigmoid(np.matmul(hidden,W2)+b2))
```

Exercise:

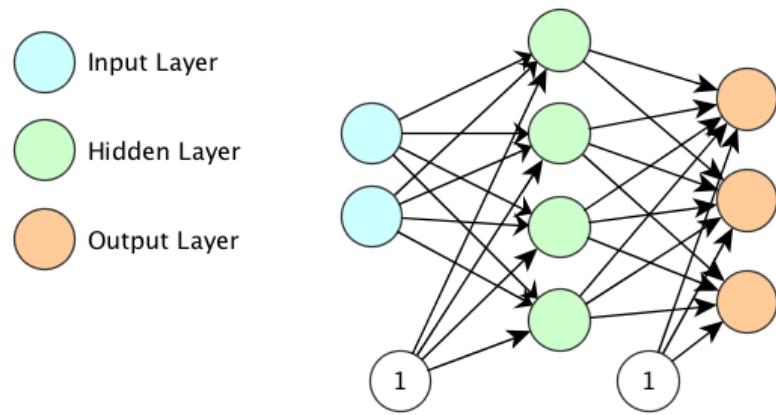


Open NB [01_simple_forward_pass.ipynb](#) and do exercise stop before Keras

Observations from NB: The benefit of hidden layers

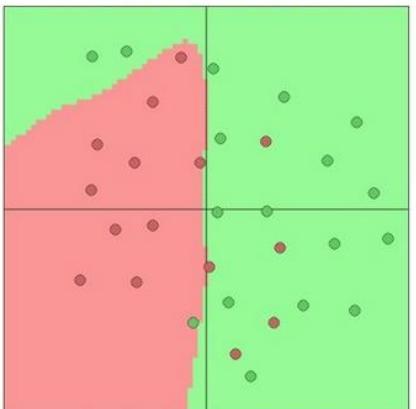


One hidden Layer

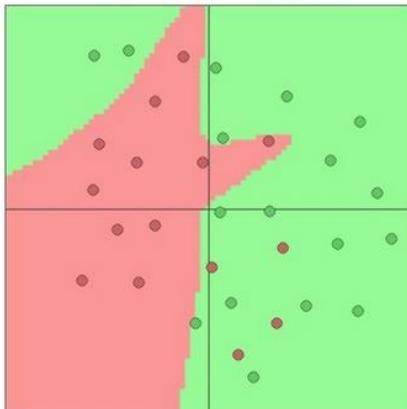


A network with one hidden layer is a universal function approximator!

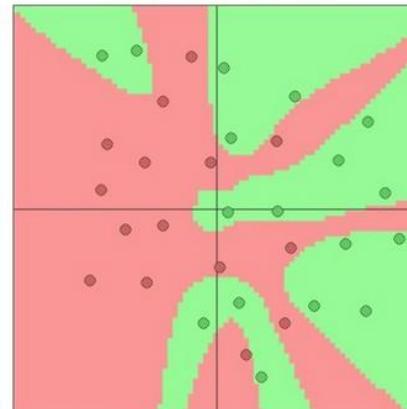
3 hidden neurons



6 hidden neurons



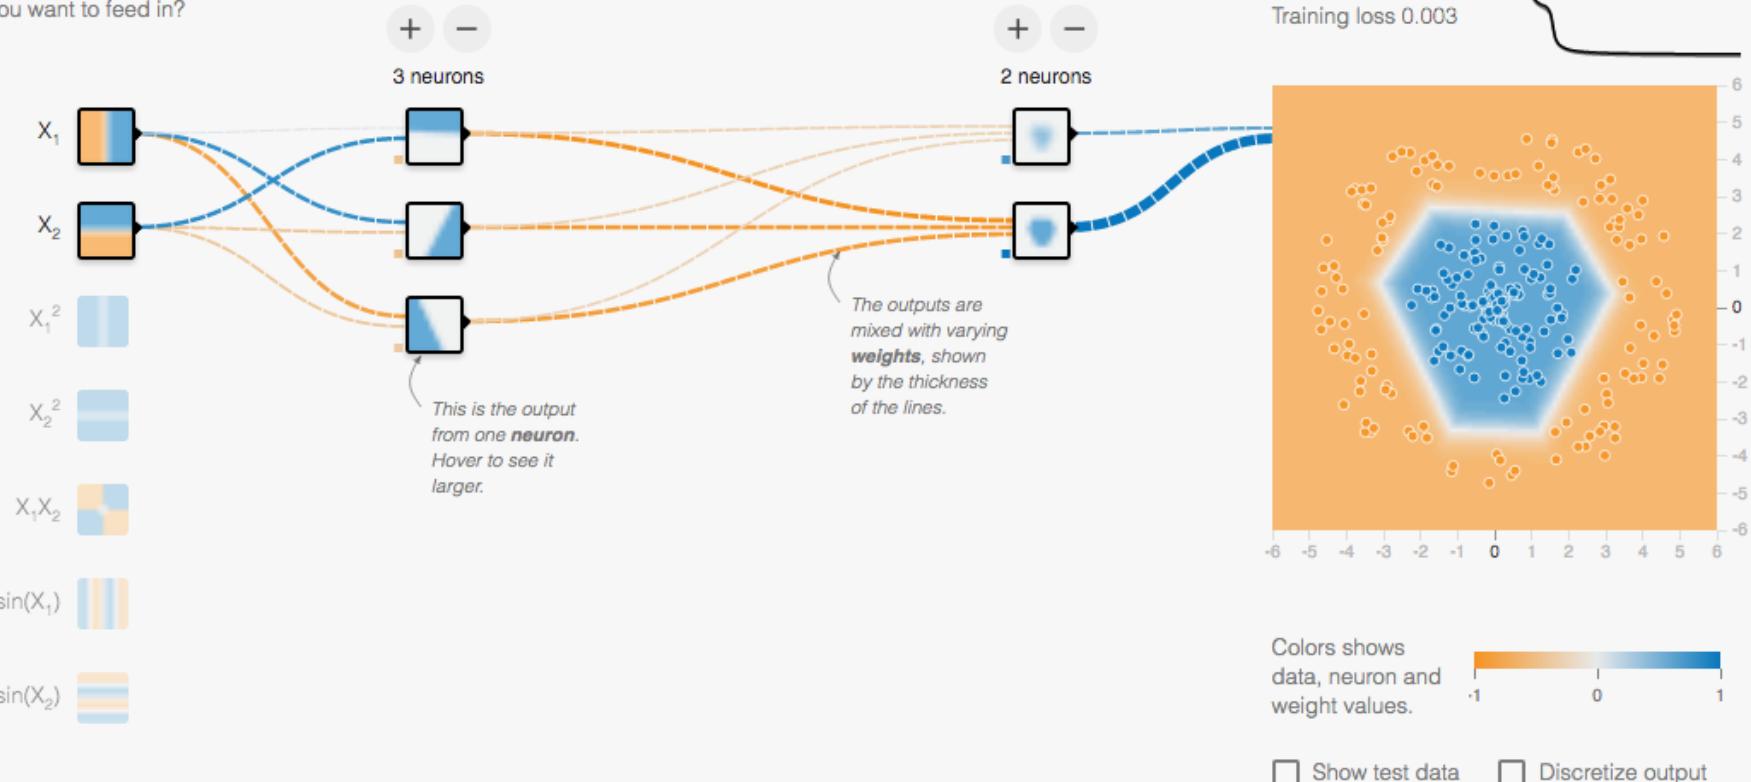
20 hidden neurons



Experiment yourself (homework)

FEATURES

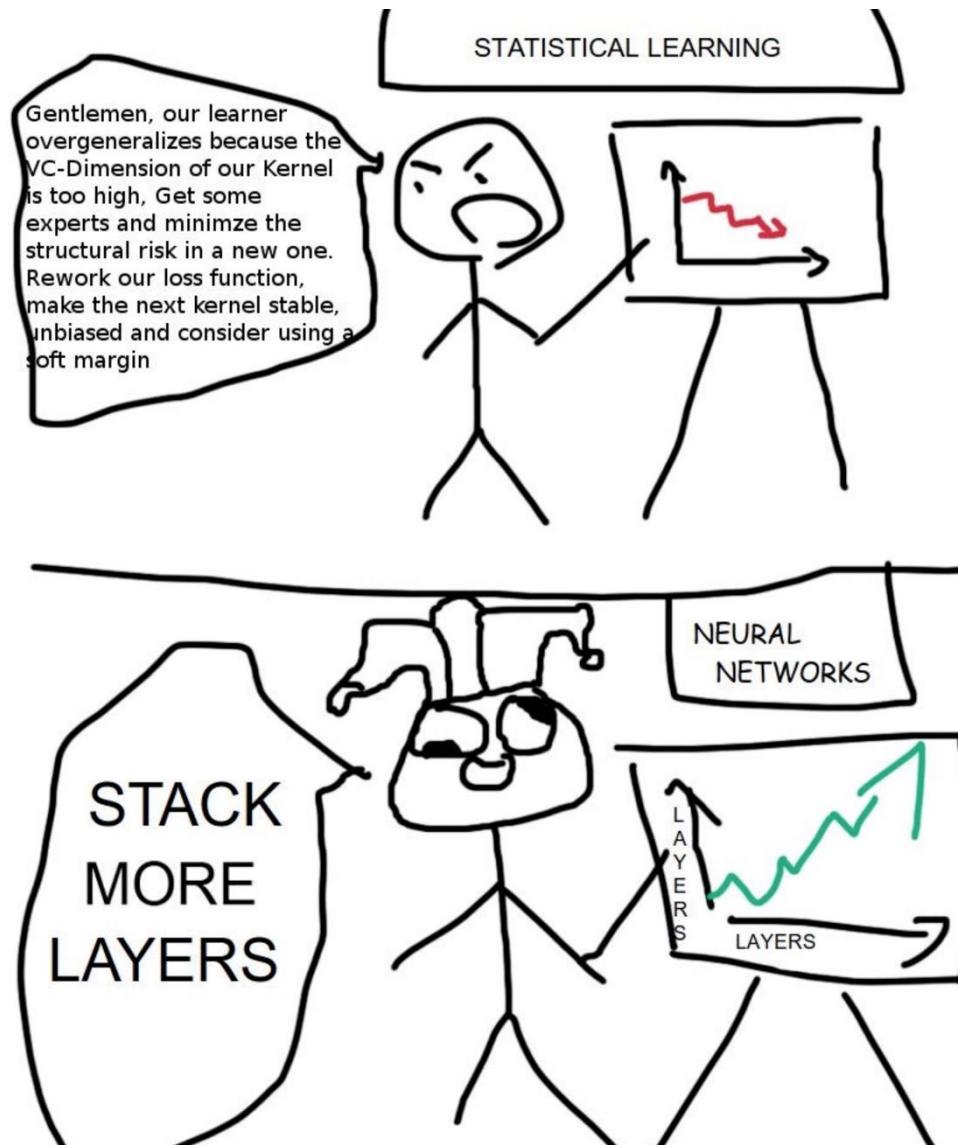
Which properties do you want to feed in?



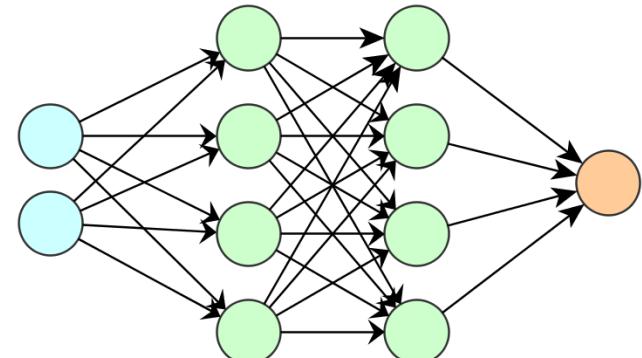
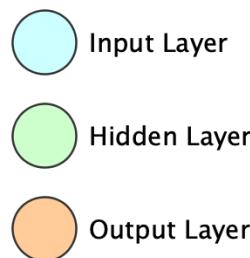
<http://playground.tensorflow.org>

Let's you explore the effect of hidden layers

DL vs Machine Learning Meme



Structure of the network



In code:

```
## Solution 2 hidden layers
def predict_hidden_2(X):
    hidden_1=sigmoid(np.matmul(X,W1)+b1)
    hidden_2=sigmoid(np.matmul(hidden_1,W2)+b2)
    return(sigmoid(np.matmul(hidden_2,W3)+b3))
```

In math ($f = \text{sigmoid}$) and $b_1=b_2=b_3=0$

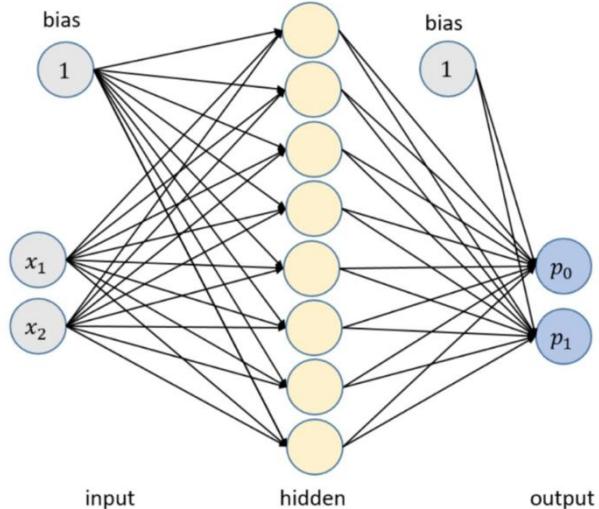
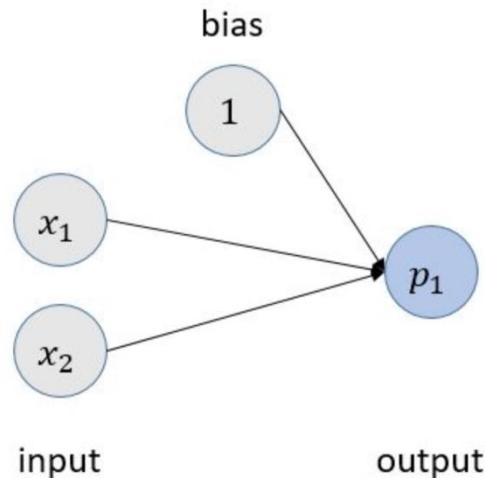
$$p = f(f(f(x W^1)W^2))$$

Looks a bit like onions, matryoshka (Russian Dolls) or lego bricks

Training NN

How to determine the weights

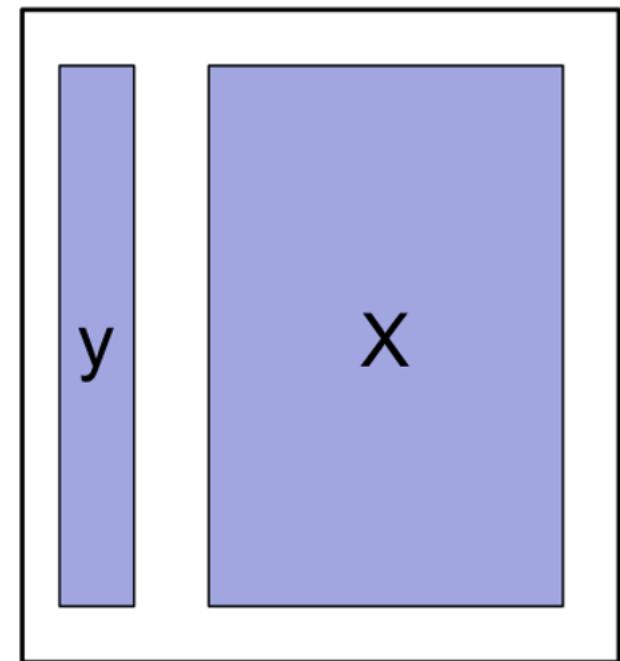
- The output of a NN is defined by the weights and biases
- These weights and biases are tuned so that the NN bests fits the training data
- The goodness of fit to the training data is quantified by a loss



Loss Functions

Tasks in DL

- The loss function depends on the task
- 2 Main tasks in DL predict y given x
 - Regression
 - Predict a number*
 - Classification
 - Predict a class*



Supervised Learning

*Later we refine this notion and predict probability distribution instead single numbers 43

Example Regression: Estimate Age From Picture

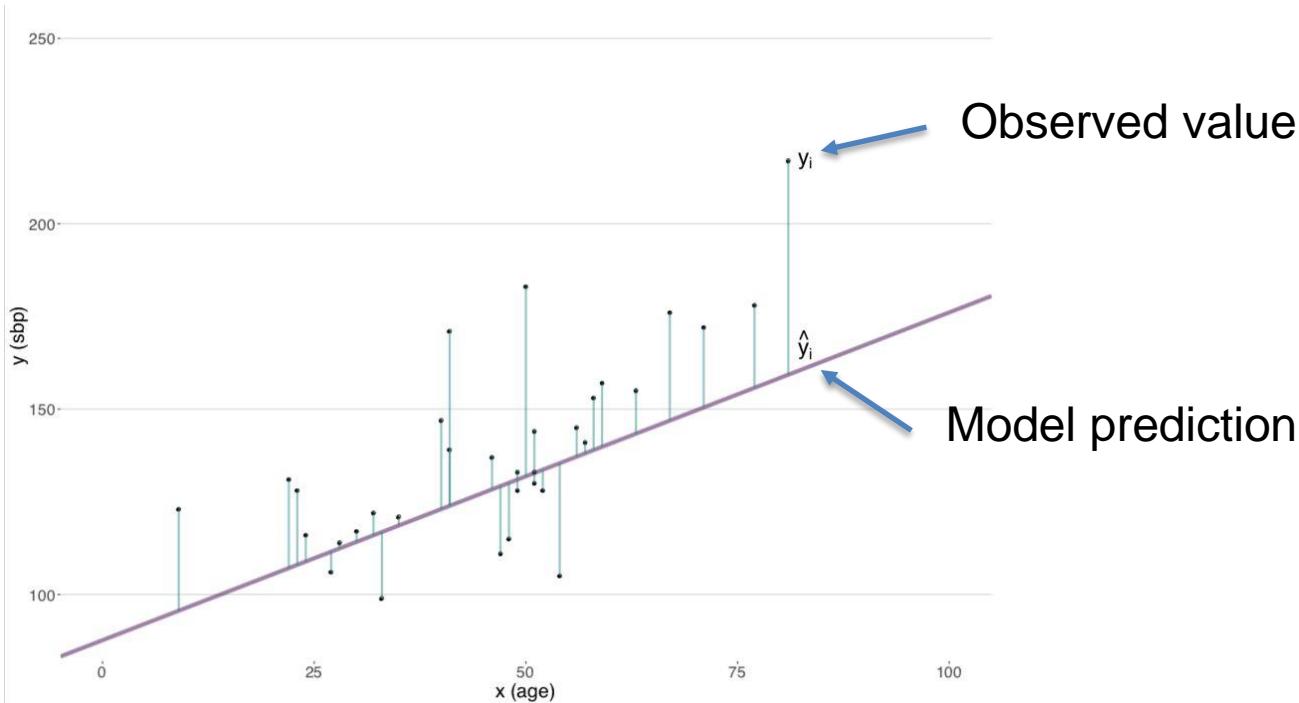


Image --> Age

<https://www.how-old.net/>

Loss Functions for Regression Problems

- Regression = Predict a number
- Example (Linear Regression)
 - Blood pressure of 31 women vs. age (training data)



Loss: Mean Squared Error (MSE) $\frac{1}{N} \sum_i (y_i - \hat{y}_i)^2$ also for regression problems (not only linear regression)

Classification

- Predict class
- Usually in DL the model predicts a probability for a class
- Example:
 - Banknote from exercise
 - Typical example Number from hand-written digit

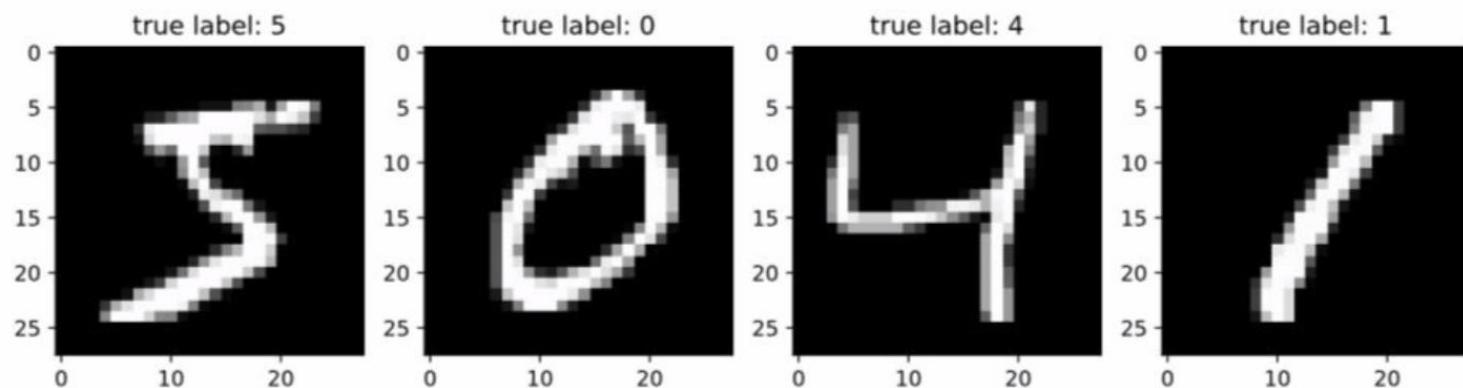
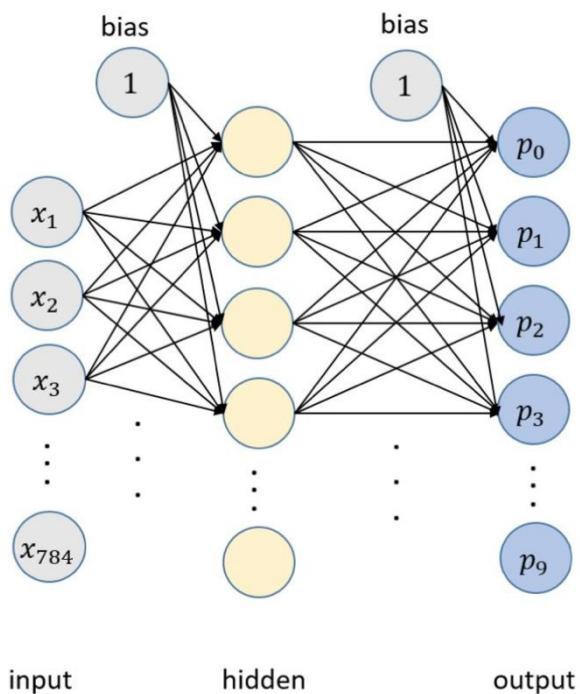


Figure 2.11 The first four digits of the MNIST data set—the standard data set used for benchmarking NN for images classification

Classification: Softmax Activation



$p_0, p_1 \dots p_9$ are probabilities for the classes 0 to 9.

Activation of last layer z_i incoming

$$p_i = \frac{e^{z_i}}{\sum_{j=0}^9 e^{z_j}}$$

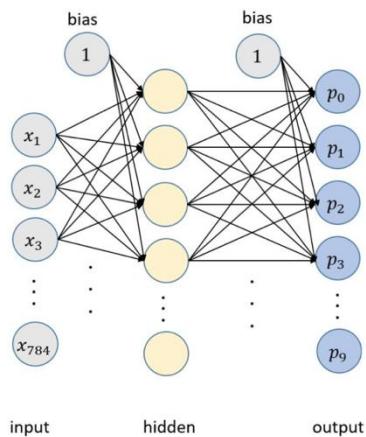
Makes outcome positive

Ensures that p_i 's sum up to one

This activation is called softmax

Figure 2.12: A fully connected NN with 2 hidden layers. For the MNIST example, the input layer has 784 values for the 28 x 28 pixels and the output layer out of 10 nodes for the 10 classes.

Loss for classification ('categorical cross-entropy')



$p_0, p_1 \dots p_9$ are probabilities for the classes 0 to 9.

- Loss is averaged of individual losses l_i of training data $i = 1, \dots N$
- Want l_i
 - 0 for perfect match, i.e. predicts class of training example $y^{(i)}$ with probability 1
 - ∞ for worst match, i.e. predicts class $y^{(i)}$ with probability 0
- $$l_i = -\log(p_{model}(y^{(i)}|x^{(i)}))$$
- $$\text{loss} = \frac{1}{N} \sum l_i$$

Training / Gradient Descent

Prinzipielle Funktionsweise: Training Bild Klassifikation

Wahre Klasse



Tiger

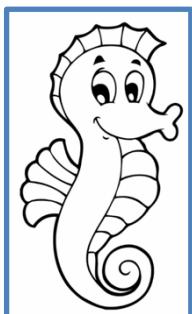
Vorhergesagtes Label

→ Seehund 🗔



Tiger

→ Tiger 🤗

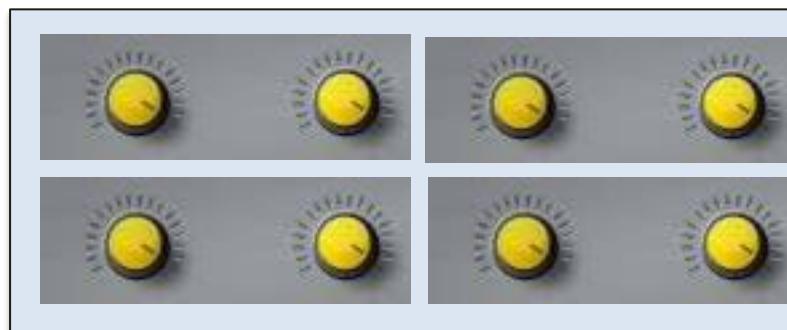


Seepferd

→ Seepferd 🤗

...

Typisch 1 Mio. Trainingsdaten



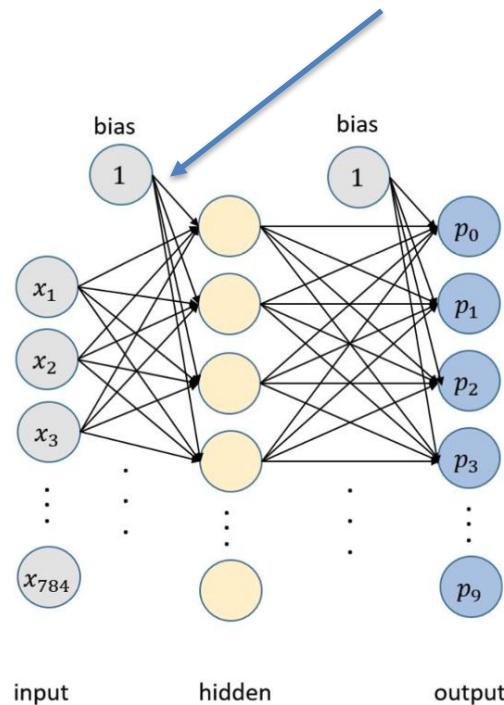
Trainingsprinzip:

Parameter werden so eingestellt, dass möglichst wenige Fehler in den Trainingsdaten gemacht werden.

Optimization in DL

- DL many parameters
 - Optimization by gradient descent
- Algorithm
 - Take a batch of training examples
 - Calculate the loss of that batch
 - Tune the parameters so that loss gets minimized

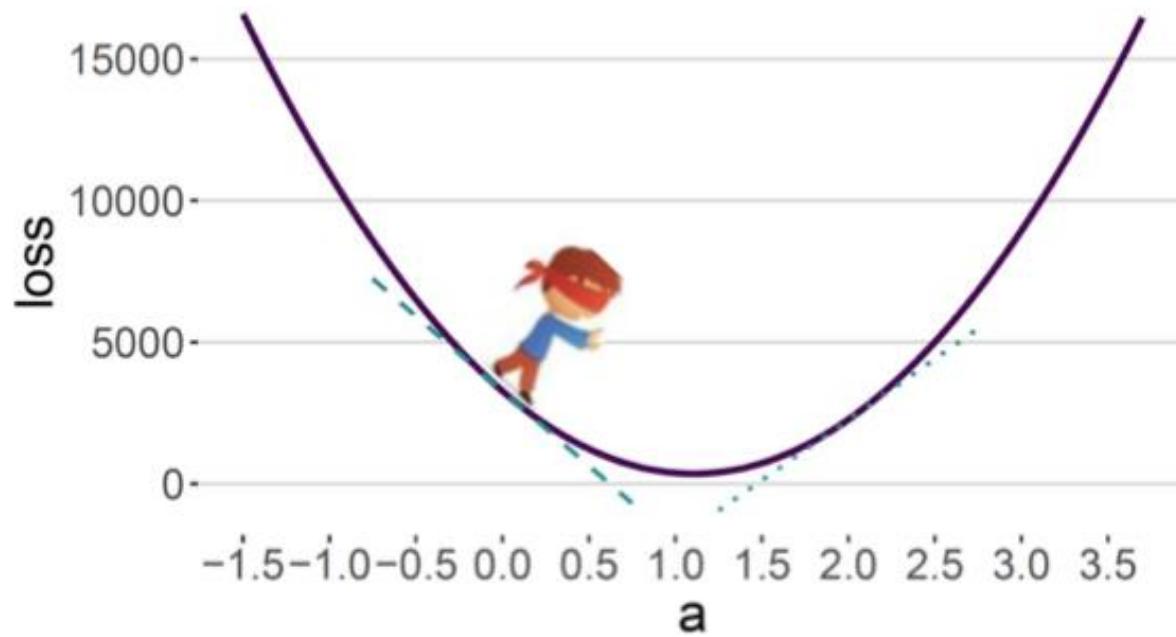
Parameters of the network are the weights.



Modern Networks have Billions (10^9) of weights. Record 2020 1.5E9
<https://openai.com/blog/better-language-models/>

Idea of gradient descent

- Shown loss function for a single parameter a



- Take a large step if slope is steep (you are away from minimum)
- Slope of loss function is given by gradient
- Iterative update of the parameters
 - $a_{t+1} = a_t - \eta \cdot \text{grad}_a(\text{loss})$

Backpropagation

- There is an efficient way to update all parameters of the network
- This is called Backpropagation (see lecture 4)
- We need to calculate the derivative of the loss function w.r.t. all weights
- Doing this efficiently (on graphic cards GPU) by hand is tedious
- Enter:
 - Deep Learning Frameworks

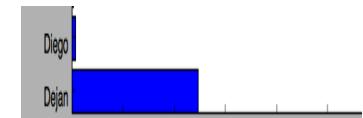
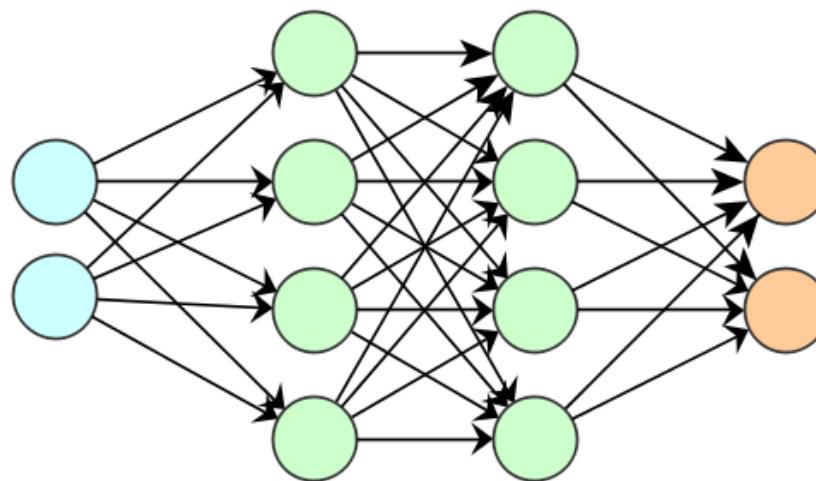
Deep Learning Frameworks

Recap: The first network

 Input Layer

 Hidden Layer

 Output Layer



- The input: e.g. intensity values of pixels of an image
 - (Almost) no pre-processing
- Information is processed layer by layer from building blocks
- Output: probability that image belongs to certain class
- Arrows are weights (these need to be learned / training)



Deep Learning Frameworks (common)

- Computation needs to be done on GPU or specialized hardware (compute performance)
- On GPU: almost exclusively on NVIDIA using the cuda library, cudnn
- Data Structure are multidimensional arrays (*tensors*) which are manipulated
- Learning require to calculate derivatives of the network w.r.t parameters.

In this course: TensorFlow with Keras

Low Level Deep Learning Libraries for Tensor Manipulations

- TensorFlow
 - Open sourced by Google Dec 2015
 - TF 1.x static computational graphs
 - Since version 2.0 also dynamic computational graphs in eager mode
- Torch / pytorch
 - Facebook, quite flexible, lua (Jan 2017 also in python, **pytorch**)
 - Dynamic computational graph
- JAX
 - New kid on the block
 - Numpy on steroids (GPU / TPU replacement)
- Chainer
 - Flexible, build graph on the fly
- MXNet
 - Can be used in many languages, build graph on the fly?
- Caffe
 - Inflexible, good of CV
 - Calculate gradients by hand
- Theano
 - Around since 2008,
 - Active development abandoned
 - Slow compiling of graph (due to optimization)



We will work (mainly) with high level libraries Keras ontop of TensorFlow

TensorFlow

What is TensorFlow

- It's API about **tensors**, which **flow** in a **computational graph**



<https://www.tensorflow.org/>

- What are **tensors**?

What is a tensor?

For TensorFlow: A tensor is an array with several indices (like in numpy). Order (a.k.a rank) are number of indices and shape is the range.

```
In [1]: import numpy as np
```

```
In [2]: T1 = np.asarray([1,2,3]) #Tensor of order 1 aka Vector  
T1
```

```
Out[2]: array([1, 2, 3])
```

```
In [3]: T2 = np.asarray([[1,2,3],[4,5,6]]) #Tensor of order 2 aka Matrix  
T2
```

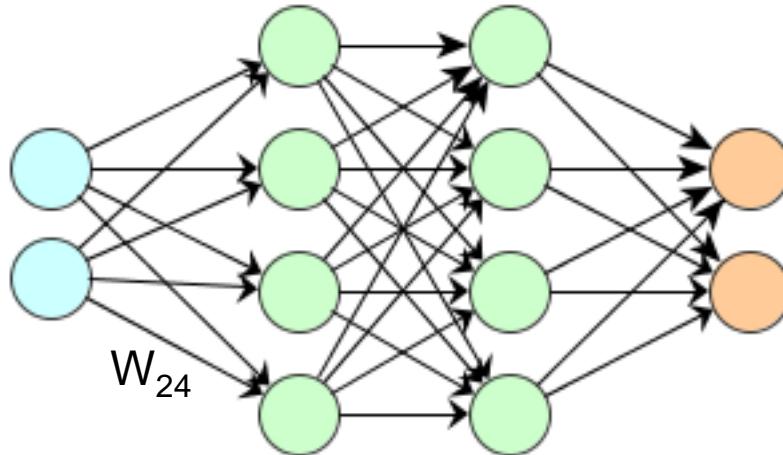
```
Out[3]: array([[1, 2, 3],  
               [4, 5, 6]])
```

```
In [4]: T3 = np.zeros((10,2,3)) #Tensor of order 3 (Volume like objects)
```

```
In [6]: print(T1.shape)  
print(T2.shape)  
print(T3.shape)
```

```
(3,)  
(2, 3)  
(10, 2, 3)
```

Typical Tensors in Deep Learning

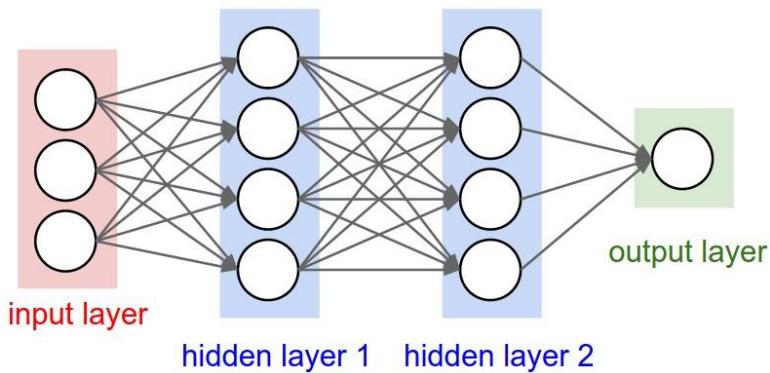


- The input can be understood as a vector
- A mini-batch of size 64 of input vectors can be understood as tensor of order 2
 - (index in batch, x_j)
- The weights going from e.g. Layer L_1 to Layer L_2 can be written as a matrix (often called W)
- A mini-batch of size 64 images with 256,256 pixels and 3 color-channels can be understood as a tensor of order 4.

Introduction to Keras

Keras as High-Level library to TensorFlow

- We use Keras as high-level library
- Libraries make use of the Lego like block structure of networks



High Level Libraries

- Keras
 - Keras is now part of TF core
 - <https://keras.io/>

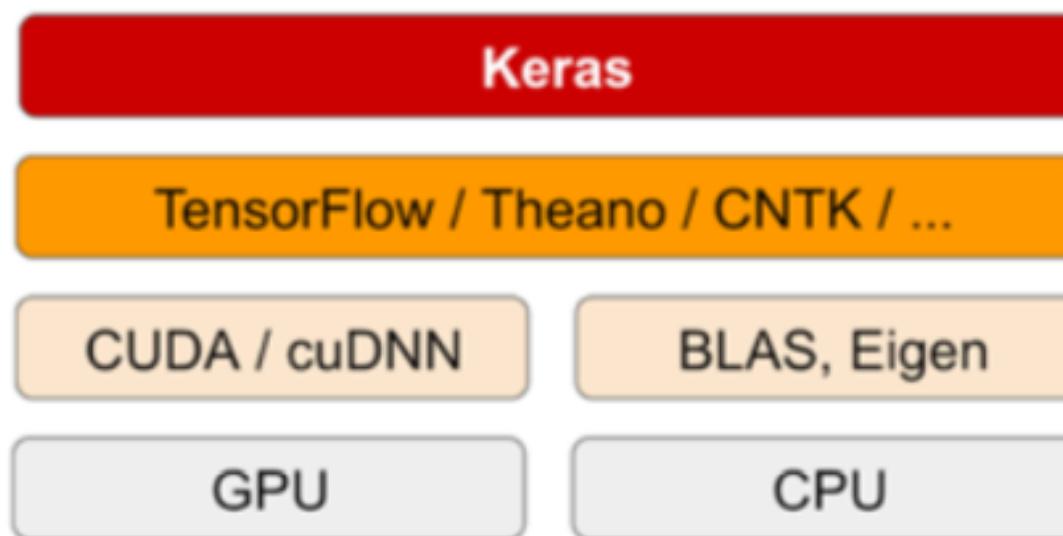
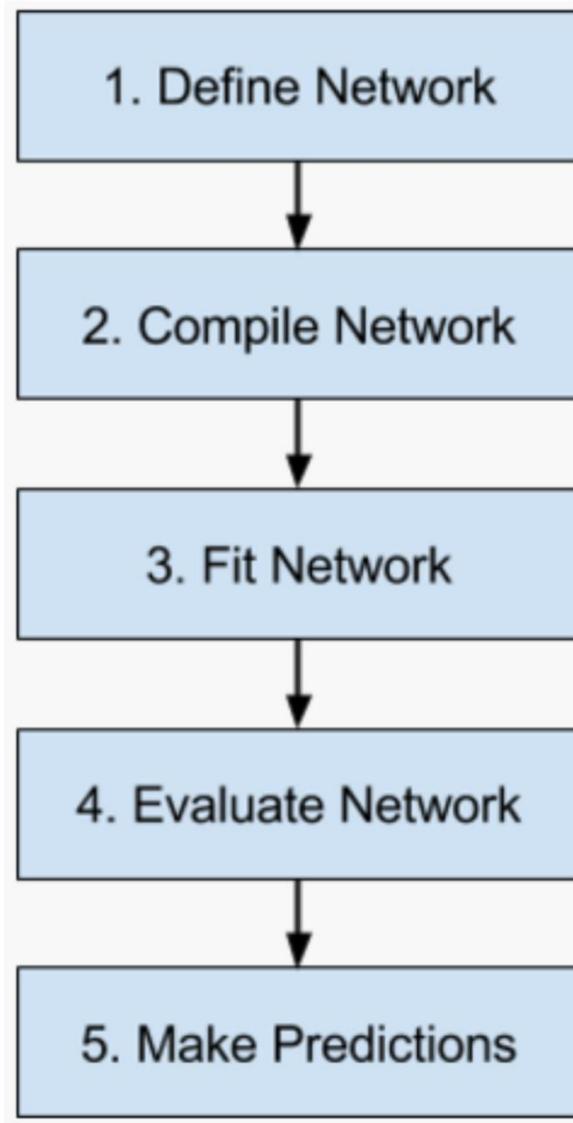


Figure: From Deep Learning with Python, Francois Chollett

Keras is multi-backend, multi-platform

- Develop in Python, R
 - On Unix, Windows, OSX
- Run the same code with...
 - TensorFlow
 - CNTK
 - Theano
 - MXNet
 - PlaidML
 - ??
- CPU, NVIDIA GPU, AMD GPU, TPU...

Keras Workflow



Define the network (layerwise)

Add loss and optimization method

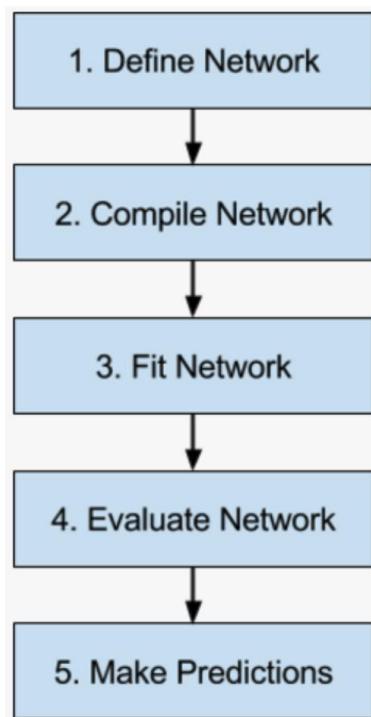
Fit network to training data

Evaluate network on test data

Use in production

A first run through

Define the network



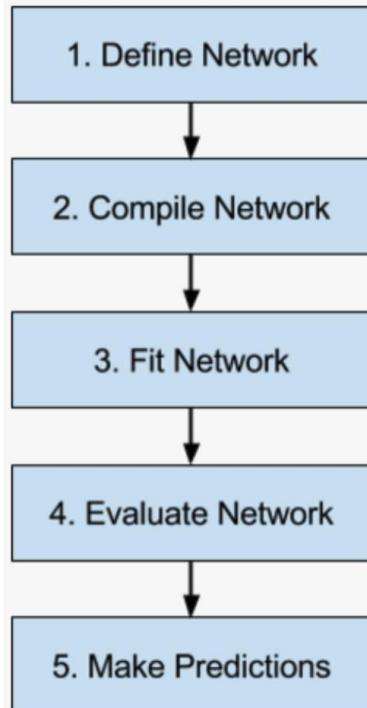
```
# define fcNN with 2 hidden layers
model = Sequential()

model.add(Dense(100, batch_input_shape=(None, 784)))
model.add(Activation('sigmoid'))
model.add(Dense(50))
model.add(Activation('sigmoid'))
model.add(Dense(10))
model.add(Activation('softmax'))
```

Input shape needs to be defined only at the beginning.

Alternative: input_dim=784

Compile the network



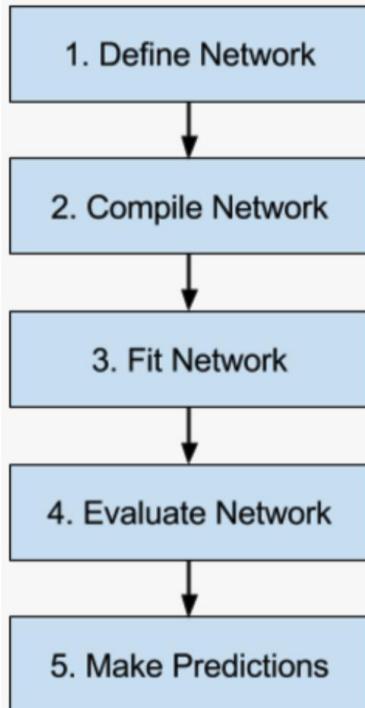
```
model.compile(loss='categorical_crossentropy',  
              optimizer='adadelta',  
              metrics=['accuracy'])
```

loss function that will be minimized

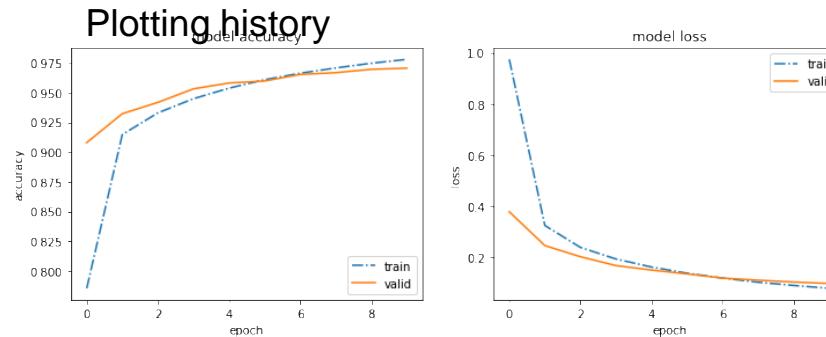
easiest optimizer is SGD
(stochastic gradient descent)

Which metrics besides 'loss' do we
want to collect during training

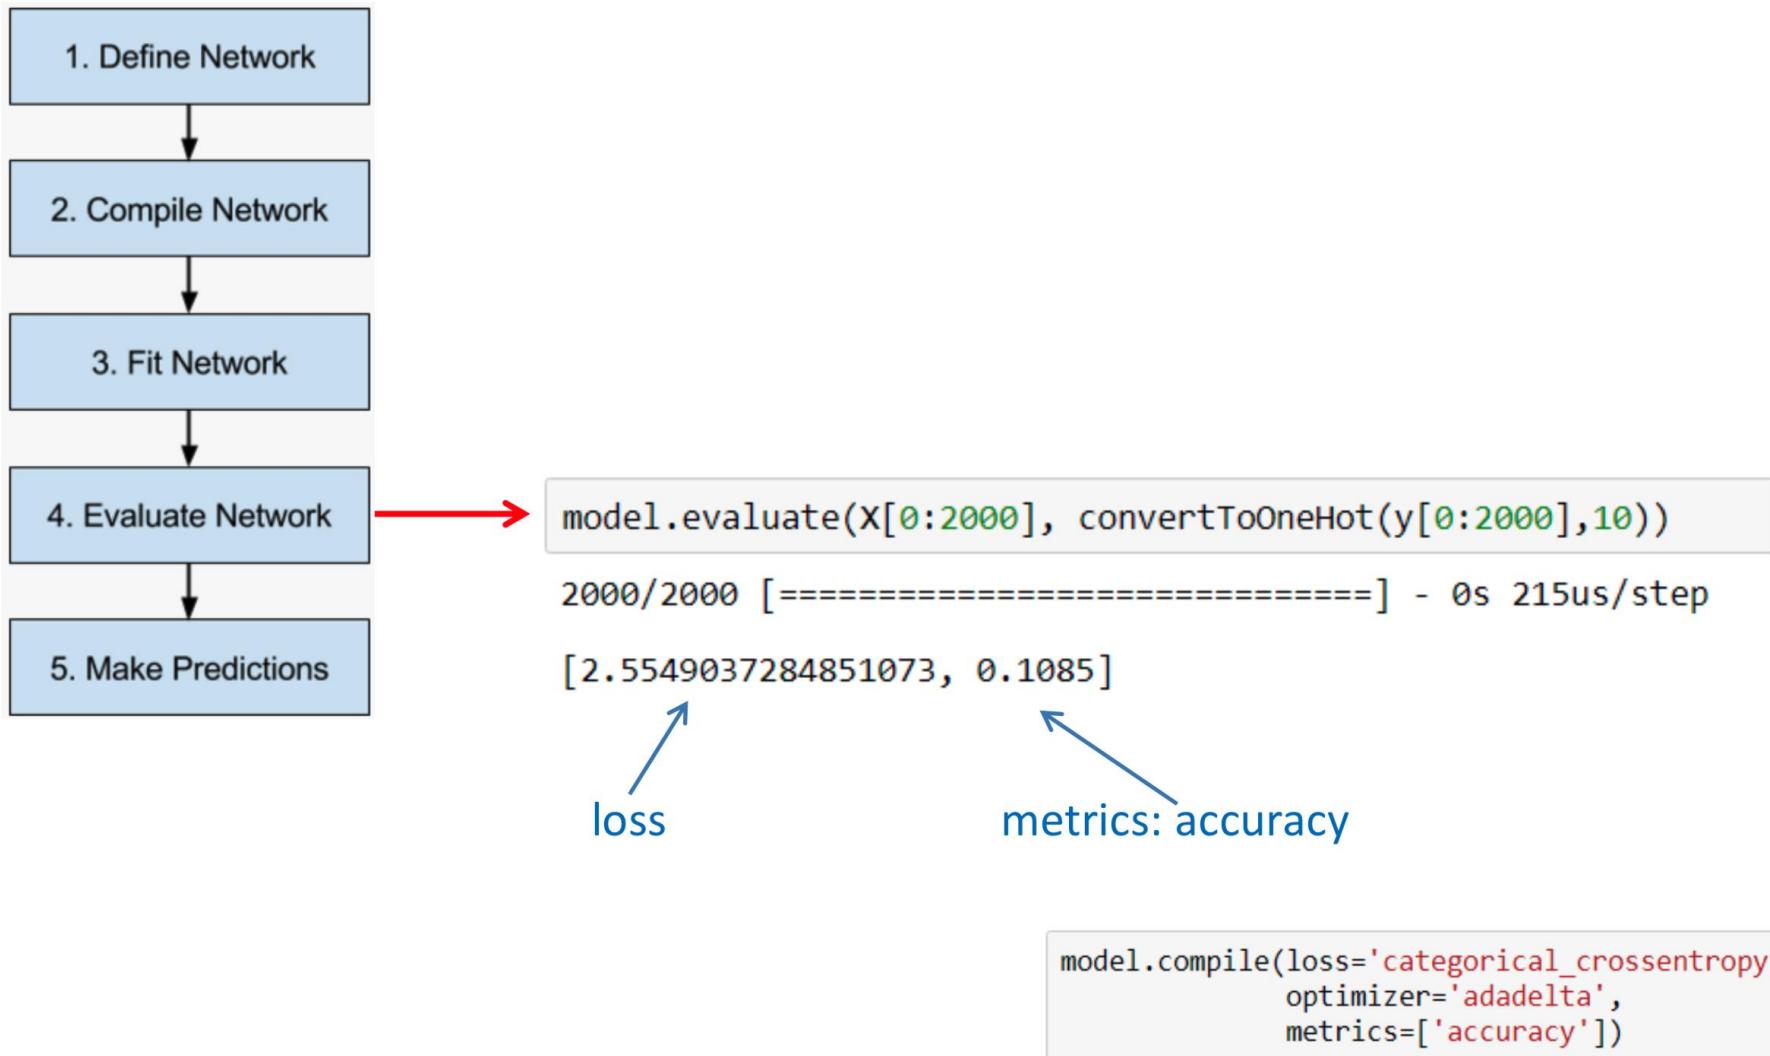
Fit the network



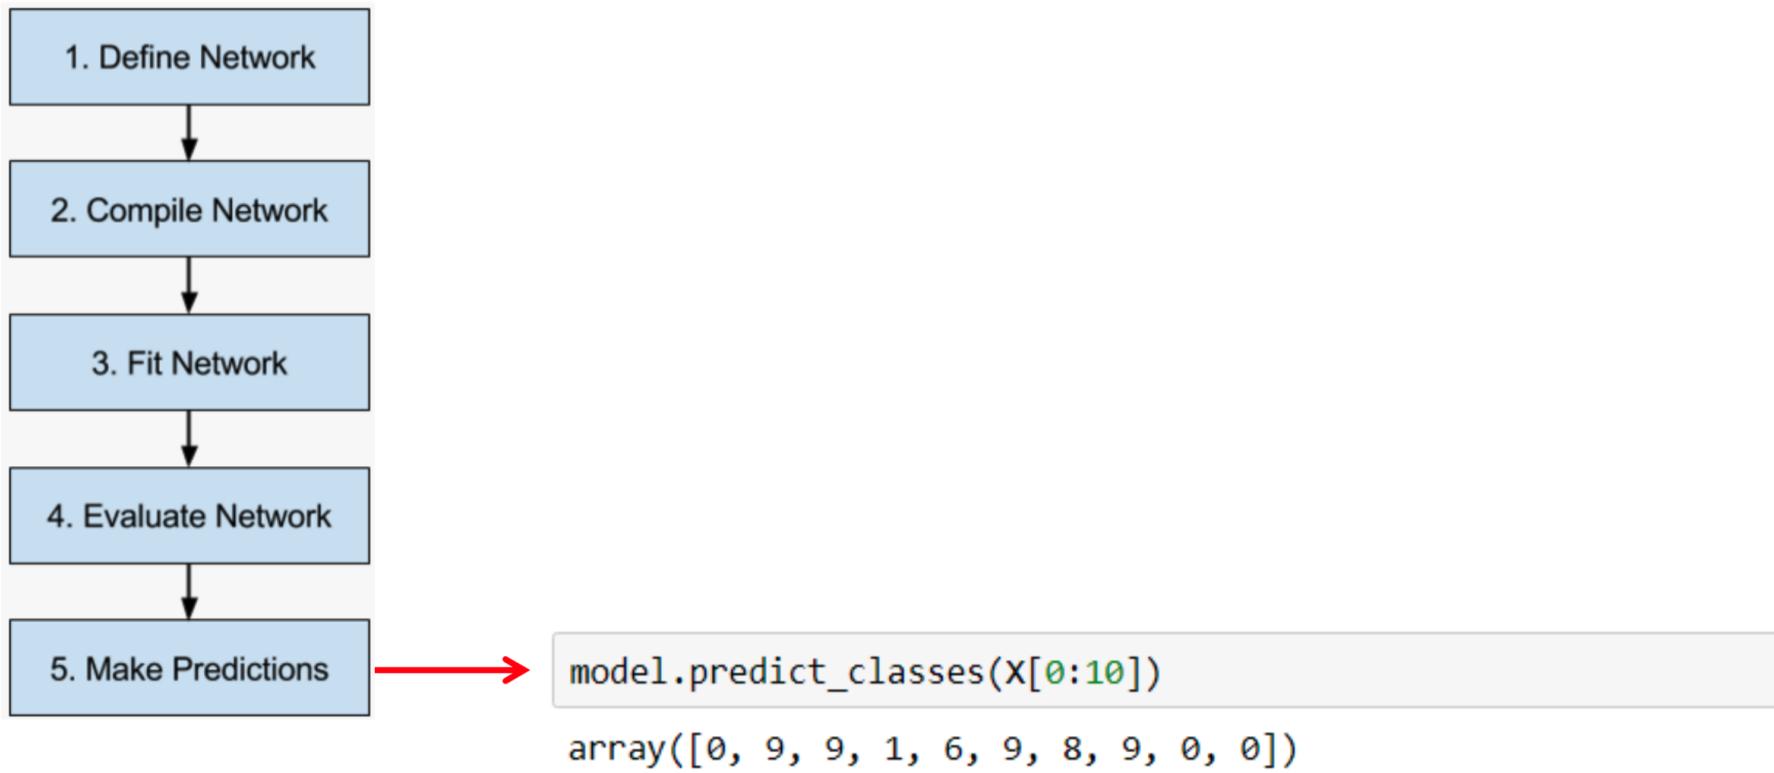
```
# train the model, memorize training history
history = model.fit(X[0:2400],  
                     convertToOneHot(y[0:2400],10),  
                     epochs=30,  
                     batch_size=128,) } input data (train)  
} output/label (train) } How and how often provide training data  
  
validation_data=[X[2400:3000],  
                 convertToOneHot(y[2400:3000],10)])
```



Evaluate the network



Make Predictions



Building NN (with keras)

- Lego Blocks (Layers)
- Way of stacking them together API Style

Building a network (API Styles)

Three API styles

- The Sequential Model
 - Dead simple
 - Only for single-input, single-output, sequential layer stacks
 - Good for 70% of use cases
- The functional API
 - Like playing with Lego bricks
 - Multi-input, multi-output, arbitrary static graph topologies
 - Good for 95% of use cases
- Model subclassing
 - Maximum flexibility
 - Larger potential error surface

Sequential API

The Sequential API

```
import keras
from keras import layers

model = keras.Sequential()
model.add(layers.Dense(20, activation='relu', input_shape=(10,)))
model.add(layers.Dense(20, activation='relu'))
model.add(layers.Dense(10, activation='softmax'))

model.fit(x, y, epochs=10, batch_size=32)
```

Functional (do you spot the error?)

The functional API

```
import keras
from keras import layers

inputs = keras.Input(shape=(10,))
x = layers.Dense(20, activation='relu')(x)
x = layers.Dense(20, activation='relu')(x)
outputs = layers.Dense(10, activation='softmax')(x)

model = keras.Model(inputs, outputs)
model.fit(x, y, epochs=10, batch_size=32)
```

Subclassing (do you spot the error?) [for completeness]

Model subclassing

```
import keras
from keras import layers

class MyModel(keras.Model):

    def __init__(self):
        super(MyModel, self).__init__()
        self.dense1 = layers.Dense(20, activation='relu')
        self.dense2 = layers.Dense(20, activation='relu')
        self.dense3 = layers.Dense(10, activation='softmax')

    def call(self, inputs):
        x = self.dense1(x)
        x = self.dense2(x)
        return self.dense3(x)

model = MyModel()
model.fit(x, y, epochs=10, batch_size=32)
```

layers



More layers

- Dropout
 - `keras.layers.Dropout`
- Convolutional (see lecture on CNN)
 - `keras.layers.Conv2D`
 - `keras.layers.Conv1D`
- Pooling (see lecture on CNN)
 - `keras.layers.MaxPooling2D`
- Recurrent (See Lecture on RNN)
 - `keras.layers.SimpleRNNCell`
 - `keras.layers.GRU`
 - `keras.layers.LSTM`
- Roll your own:
 - Implement `keras.layers.Layer` class
 - <https://keras.io/layers/writing-your-own-keras-layers/>



Activation

```
keras.layers.Activation(activation)
```

Applies an activation function to an output.

Arguments

e.g. 'relu', 'softmax', 'tanh', ...

- **activation:** name of activation function to use (see: [activations](#)), or alternatively, a Theano or TensorFlow operation.

```
from keras.layers import Activation, Dense  
  
model.add(Dense(64))  
model.add(Activation('tanh'))
```

This is equivalent to:

```
model.add(Dense(64, activation='tanh'))
```

Note: Activations are also layers

Output Layer

The last layer of the network is a bit special

- Classification
 - # of nodes = # of classes
 - For binary classification sometime only probability of class 1 is reported
 - Usually output is probability for class
 - Use softmax in that case
- Regression (simple distributions)
 - In the interpretation the output of a NN is a probability
 - #nodes = 1
 - Gaussian with fixed variance (the usual regression)
 - Poisson (count data) see later
- Regression (more complicated distributions)
 - #nodes = #number of parameters for distribution

Tasks



1. Finish NB 01_simple_forward_pass.ipynb have a look at the Keras implementation
2. Do exercise NB 02

https://github.com/tensorchiefs/dl_course_2023/blob/master/notebooks/02_fcnn_with_banknote.ipynb